

HAMMLET: AN INFINITE HIDDEN MARKOV MODEL
WITH LOCAL TRANSITIONS

by
Colin Reimer Dawson

A Dissertation Submitted to the Faculty of the

PROGRAM IN STATISTICS

In Partial Fulfillment of the Requirements

For the Degree of

DOCTOR OF PHILOSOPHY

In the Graduate College

THE UNIVERSITY OF ARIZONA

2 0 1 6

THE UNIVERSITY OF ARIZONA
GRADUATE COLLEGE

As members of the Dissertation Committee, we have read the dissertation
prepared by Colin Reimer Dawson
entitled “HaMMLeT: An Infinite Hidden Markov Model with Local Transitions”
and recommend that it be accepted as fulfilling the dissertation requirement for the
Degree of Doctor of Philosophy.

_____ Date: 08/08/2016
Clayton Morrison

_____ Date: 08/08/2016
Katherine Barnes

_____ Date: 08/08/2016
Helen Zhang

Final approval and acceptance of this dissertation is contingent upon the candidate’s
submission of the final copies of the dissertation to the Graduate College.

I hereby certify that I have read this dissertation prepared under my direction and
recommend that it be accepted as fulfilling the dissertation requirement.

_____ Date: 08/08/2016
Dissertation Director: Clayton Morrison

STATEMENT BY AUTHOR

This dissertation has been submitted in partial fulfillment of requirements for an advanced degree at The University of Arizona and is deposited in the University Library to be made available to borrowers under rules of the Library.

Brief quotations from this dissertation are allowable without special permission, provided that accurate acknowledgment of source is made. Requests for permission for extended quotation from or reproduction of this manuscript in whole or in part may be granted by the head of the major department or the Dean of the Graduate College when in his or her judgment the proposed use of the material is in the interests of scholarship. In all other instances, however, permission must be obtained from the author.

SIGNED: _____

ACKNOWLEDGEMENTS

I would like to acknowledge the following people for their support, intellectual, moral, emotional, and otherwise.

- Clay Morrison, my advisor, for his tireless support, for helping me stay organized, and for challenging me to be able to explain myself.
- Kobus Barnard, for turning me on to Bayesian Modeling in the first place, way back in the aughts sometime, and for providing the resources and opportunities for me to actually learn to put ideas to code.
- The rest of my committee, Kathie Barnes and Helen Zhang, for their valuable feedback.
- Paul Cohen, for bringing me in to SISTA, providing the opportunity (and funding!) to work on so many varied challenging research projects, and giving me so much leeway in teaching for the information science program.
- Kyle Simek, Ernesto Brau, and Andrew Prehoehl, for being excellent lab mates, for putting together the summer “coding bootcamps” that were so valuable to me as someone without a lot of formal CS training, for the very stimulating reading groups, and for being there to bounce ideas off of.
- My first research student at Oberlin, Bill Huang, for writing code behind several of the figures herein.
- My parents, for being able to smile and nod when I said I was going to do *another* Ph.D., and for not asking “have you started writing your dissertation yet?” more than once a week.
- My children, Arlo and Esai, for being there with snuggles and laughter every day, and reminding me what’s actually important.
- My wife, Sarah, for everything and more. There are no words.

DEDICATION

To Sarah, my amazing wife, for putting up with me during not one, but *two* rounds of graduate school. Next, the tenure track. What fun for her!

TABLE OF CONTENTS

LIST OF TABLES	9
LIST OF FIGURES	10
ABSTRACT	11
1. STATISTICAL MODELS AND MODEL SELECTION	14
1.1. Introduction	14
1.1.1. An overview of the dissertation	16
1.2. Model Selection	17
1.2.1. The Likelihood Function and the Conservation of Explanatory Power	17
1.2.2. The Bias-Variance Tradeoff	19
1.2.3. Example: Polynomial Regression	20
1.2.4. Balancing Fit and Complexity: Bayesian Occam's Razor	22
1.3. Clustering via a Mixture of Gaussians Model	24
1.4. Clustering Sequential Data: The Hidden Markov Model	25
1.5. Model Selection in HMMs and the Mixture of Gaussians Model	27
2. BAYESIAN NONPARAMETRIC MODELS AND THE INFINITE HIDDEN MARKOV MODEL	29
2.1. Parametric Vs. Nonparametric Models	29
2.2. The Dirichlet Process	30
2.2.1. The Normalized Gamma Process representation of the DP	32
2.2.2. The Stick-Breaking Process construction of the DP	33
2.2.3. The Chinese Restaurant Process	33
2.2.4. The Dirichlet Process Mixture Model	34
2.2.5. A Gibbs Sampler for DP Mixture Models	35
2.3. An Infinite State HMM	36
2.3.1. The Hierarchical Dirichlet Process	36
2.3.2. The HDP-HMM	36
2.4. Inference in Finite and Infinite State HMMs	36
2.4.1. The Forward Backward Algorithm	36
2.4.2. Two Gibbs Samplers for the HDP-HMM	36
2.4.3. Beam Sampling in the HDP-HMM	36

TABLE OF CONTENTS—*Continued*

3. HAMMLET: AN INFINITE HIDDEN MARKOV MODEL WITH LOCAL TRANSITIONS	37
3.1. Transition Dynamics in the HDP-HMM	38
3.2. An HDP-HMM With Local Transitions	40
3.2.1. A Normalized Gamma Process representation of the HDP-HMM	41
3.2.2. Promoting “Local” Transitions	42
3.3. The HDP-HMM-LT as a continuous-time Markov Jump Process with “failed” jumps	44
3.3.1. An HDP-HSMM-LT modification	47
3.3.2. Summary	49
3.4. MCMC Inference in the “Failed Jumps” Representation	49
3.4.1. Sampling π , β , α and γ	50
3.4.2. Sampling z and the auxiliary variables	55
3.4.3. Sampling state and emission parameters	56
3.5. Use Cases	57
4. BINARY VECTOR STATES: SPEAKER DIARIZATION	59
4.1. Binary State Vectors	59
4.1.1. Additional Inference Steps	60
4.1.2. Summary	63
4.2. “Cocktail Party” Data	65
4.3. Synthetic Data Without Local Transitions	67
5. CATEGORICAL VECTOR STATES: POWER DISAGGREGATION	70
5.1. Generalizing to Categorical-Valued θ	70
5.2. Priors and Representations in the Categorical State Variant	71
5.3. Adapting Posterior Inference for Categorical State Vectors	72
5.4. Power Disaggregation	73
6. SEPARATE SIMILARITIES AND EMISSIONS: LEARNING TONAL GRAMMAR IN MUSIC	76
6.1. Separable Similarity and Emissions	76
6.2. A Hamiltonian Monte Carlo step to sample η	76
6.3. Synthetic Data from an HMM with a Nearly Block Diagonal Transition Matrix	79
6.4. Discovering Chord Equivalence Classes in Tonal Music	80
7. CONCLUSIONS AND FUTURE WORK	81

TABLE OF CONTENTS—*Continued*

REFERENCES	82
----------------------	----

LIST OF TABLES

LIST OF FIGURES

FIGURE 1.1. Polynomials of varying degrees fit using ordinary least squares (equivalently, maximum likelihood estimation) to a set of eleven points generated from a fifth order polynomial with independent Normal residuals. 1.1a–1.1c: fits of polynomial order 1, 2 and 10. Filled circles are training data, unfilled circles are data not used during fitting. 1.1d squared prediction error in and out of the training sample for polynomial orders 1 through 10.	21
FIGURE 4.1. (a-b) Accuracy and F1 scores for the HDP-HMM-LT, standard HDP-HMM, and Binary Factorial HMM on the Cocktail Party Data. Metrics are averaged over 10 Gibbs runs on each model, with error bars representing a 99% confidence interval for the mean per iteration. The first 100 iterations are not shown. (c) Learned similarity parameter, λ , for the LT model, (d) Number of distinct states used by HDP-HMM and HDP-HMM-LT. The first 100 iterations are excluded.	68
FIGURE 4.2. (a-b) Accuracy and F1 for the three models on data generated from an HDP-HMM without local transitions, (c) Learned similarity parameter, λ , for the LT model, (d) Number of states used by the HDP-HMM and HDP-HMM-LT. The first 100 iterations are omitted.	69
FIGURE 5.1. A sample data interval from the REDD dataset Kolter and Johnson (2011). The top channel contains the total measured power in watts consumed by a home during a period of approximately 24 hours. Each timestep represents a 20 second intervals, during which the amplitude recorded is a median of the amplitudes in the original higher resolution data.	75

ABSTRACT

In a classical mixture modeling, each data point is modeled as arising i.i.d. (typically) from a weighted sum of probability distributions, where both the weights and the parameters of the mixture components are targets of inference. When data arises from different sources that may not give rise to the same mixture distribution, a hierarchical model can allow the source contexts to share components while assigning different weights across them (while perhaps coupling the weights to “borrow strength” across contexts). The Dirichlet Process (DP) Mixture Model (e.g., Rasmussen (2000)) is a Bayesian approach to mixture modeling which models the data as arising from a countably infinite number of components: the Dirichlet Process provides a prior on the mixture weights that guards against overfitting. The Hierarchical Dirichlet Process (HDP) Mixture Model (Teh et al., 2006) employs a separate DP Mixture Model for each context, but couples the weights across contexts by using a common base measure which is itself drawn from a top-level DP. This coupling is critical to ensure that mixture components are reused across contexts. For example, in natural language topic modeling, a common application domain for mixture models, the components represent semantic topics, and the contexts are documents, and it is critical that topics be reused across documents.

These models have been widely adopted in Bayesian statistics and machine learning. However, a limitation of DPs is that the atoms are *a priori* exchangeable, and in the case of HDPs, the component weights are independent conditioned on the top-level measure. This is unrealistic in many applications, including topic modeling, where certain components (e.g., topics) are expected to correlate across contexts (e.g., documents). In the case of topic modeling, the Discrete Infinite Logistic Normal model (DILN; Paisley et al. (2011)) addresses this shortcoming by associating with each mixture component a latent location in an abstract metric, and rescaling each

context-specific set of weights, initially drawn from an HDP, by an exponentiated draw from a Gaussian Process (GP), so that components which are nearby in space tend to have their weights be scaled up or down together. However, inference in this model requires the posterior distribution to be approximated by a variational family, as MCMC sampling from the exact posterior was deemed intractable. Thus, one goal of this dissertation is the development of simple MCMC algorithms for correlated components.

A second application of HDPs is to time series models, in particular Hidden Markov Models (HMMs), where the HDP can be used as a prior on a doubly infinite transition matrix for the latent Markov chain, giving rise to the HDP-HMM (first developed, as the “Infinite HMM”, by Beal et al. (2001), and subsequently shown to be a case of an HDP by Teh et al. (2006)). There, the hierarchy is over rows of the transition matrix, and the distributions across rows are coupled through a top-level Dirichlet Process. The sequential nature of the problem introduces two added wrinkles, namely that: the contexts themselves are random (since the context when generating state t is the state at time $t - 1$), and the set of contexts is the same as the set of components. Hence, not only might the components be correlated with each other via locations in some latent space, but we might expect that contexts that correspond to correlated components will overall have similar distributions.

In the first part of the dissertation, I will present a formal overview of Mixture Models and Hidden Markov Models. I then turn to a discussion of Dirichlet Processes and their various representations, as well as associated schemes for tackling the problem of doing approximate inference over an infinitely flexible model with finite computational resources. I will then turn to the Hierarchical Dirichlet Process (HDP) and its application to an infinite state Hidden Markov Model, the HDP-HMM.

The central contribution of the dissertation is a novel probabilistic model, which I call the Hierarchical Dirichlet Process Hidden Markov Model With Local Transitions (HDP-HMM-LT, or HaMMLLeT for short), which achieves the goal of simultaneously

modeling correlations between contexts and components by assigning each a location in a metric space and promoting transitions between states that are near each other. I present a Gibbs sampling scheme for inference in this model, employing an augmented data representation to simplify the relevant conditional distributions. I give a intuitive interpretation of the augmented representation by casting the discrete time chain as a continuous time chain in which durations are not observed, and in which some jump attempts fail and are never observed. By tying the success probability of a jump between two states to the distance between them, the first successful (and therefore observed) jump is more likely to be to a nearby state. I refer to this representation as a Markov Process With Failed Jumps. I test this model on several synthetic and real data sets, showing that in at least some circumstances, the HaMMLeT model more effectively finds the latent time series structure underlying the observations.

HAMMLET: AN INFINITE HIDDEN MARKOV MODEL WITH LOCAL TRANSITIONS

Colin Reimer Dawson, Ph.D.
The University of Arizona, 2016

Director: Clayton Morrison

In a classical mixture modeling, each data point is modeled as arising i.i.d. (typically) from a weighted sum of probability distributions, where both the weights and the parameters of the mixture components are targets of inference. When data arises from different sources that may not give rise to the same mixture distribution, a hierarchical model can allow the source contexts to share components while assigning different weights across them (while perhaps coupling the weights to “borrow strength” across contexts). The Dirichlet Process (DP) Mixture Model (e.g., Rasmussen (2000)) is a Bayesian approach to mixture modeling which models the data as arising from a countably infinite number of components: the Dirichlet Process provides a prior on the mixture weights that guards against overfitting. The Hierarchical Dirichlet Process (HDP) Mixture Model (Teh et al., 2006) employs a separate DP Mixture Model for each context, but couples the weights across contexts by using a common base measure which is itself drawn from a top-level DP. This coupling is critical to ensure that mixture components are reused across contexts. For example, in natural language topic modeling, a common application domain for mixture models, the components represent semantic topics, and the contexts are documents, and it is critical that topics be reused across documents.

These models have been widely adopted in Bayesian statistics and machine learning. However, a limitation of DPs is that the atoms are *a priori* exchangeable, and in the case of HDPs, the component weights are independent conditioned on the top-level measure. This is unrealistic in many applications, including topic modeling, where certain components (e.g., topics) are expected to correlate across contexts

(e.g., documents). In the case of topic modeling, the Discrete Infinite Logistic Normal model (DILN; Paisley et al. (2011)) addresses this shortcoming by associating with each mixture component a latent location in an abstract metric, and rescaling each context-specific set of weights, initially drawn from an HDP, by an exponentiated draw from a Gaussian Process (GP), so that components which are nearby in space tend to have their weights be scaled up or down together. However, inference in this model requires the posterior distribution to be approximated by a variational family, as MCMC sampling from the exact posterior was deemed intractable. Thus, one goal of this dissertation is the development of simple MCMC algorithms for correlated components.

A second application of HDPs is to time series models, in particular Hidden Markov Models (HMMs), where the HDP can be used as a prior on a doubly infinite transition matrix for the latent Markov chain, giving rise to the HDP-HMM (first developed, as the “Infinite HMM”, by Beal et al. (2001), and subsequently shown to be a case of an HDP by Teh et al. (2006)). There, the hierarchy is over rows of the transition matrix, and the distributions across rows are coupled through a top-level Dirichlet Process. The sequential nature of the problem introduces two added wrinkles, namely that: the contexts themselves are random (since the context when generating state t is the state at time $t - 1$), and the set of contexts is the same as the set of components. Hence, not only might the components be correlated with each other via locations in some latent space, but we might expect that contexts that correspond to correlated components will overall have similar distributions.

In the first part of the dissertation, I will present a formal overview of Mixture Models and Hidden Markov Models. I then turn to a discussion of Dirichlet Processes and their various representations, as well as associated schemes for tackling the problem of doing approximate inference over an infinitely flexible model with finite computational resources. I will then turn to the Hierarchical Dirichlet Process (HDP) and its application to an infinite state Hidden Markov Model, the HDP-HMM.

The central contribution of the dissertation is a novel probabilistic model, which I call the Hierarchical Dirichlet Process Hidden Markov Model With Local Transitions (HDP-HMM-LT, or HaMMLeT for short), which achieves the goal of simultaneously modeling correlations between contexts and components by assigning each a location in a metric space and promoting transitions between states that are near each other. I present a Gibbs sampling scheme for inference in this model, employing an augmented data representation to simplify the relevant conditional distributions. I give an intuitive interpretation of the augmented representation by casting the discrete time chain as a continuous time chain in which durations are not observed, and in which some jump attempts fail and are never observed. By tying the success probability of a jump between two states to the distance between them, the first successful (and therefore observed) jump is more likely to be to a nearby state. I refer to this representation as a Markov Process With Failed Jumps. I test this model on several synthetic and real data sets, showing that in at least some circumstances, the HaMMLeT model more effectively finds the latent time series structure underlying the observations.

1. STATISTICAL MODELS AND MODEL SELECTION

1.1. Introduction

The statistician George Box famously said that “Essentially, all models are wrong, but some are useful.” (Box et al., 1987). That is, by necessity, a statistical model is a simplification of the phenomenon that it is modeling, and indeed, increasing fidelity is not always a good thing. For one, more complex models are more difficult to understand, which can be important if the goal of modeling is to achieve some insight about the law-like behavior governing a natural phenomenon. But even if the model is being applied strictly for the purposes of making predictions, and being able to gain some verbal understanding is not a priority, the more complex a model becomes, the more sensitive it is to the idiosyncracies of the data that it is informed by.

An important question in statistics is how to go about choosing a model with enough complexity and flexibility to account for the interesting regularities in the process that produced the data, but not so much complexity that the model will not generalize; after all, the purpose of statistical modeling is almost always to explain something and/or to make predictions about the future, not simply to describe data already collected.

A number of techniques have been developed in both frequentist and Bayesian statistics to select from among a set of model forms one that balances fit and generalizability, that is, to do **model selection**. An alternative to model selection is to employ a model form whose complexity is not fixed, but instead adapts to become increasingly flexible as the data set grows larger: this is the **nonparametric** approach, which has a long history in frequentist statistics (e.g., Rosenblatt et al. (1956); Parzen (1962)), where the typical approach is to model the data without any assumptions about the specific distributional family that generated the observations.

Distribution-free estimators have been shown to have good asymptotic properties (see, e.g., Wasserman (2007) for a review of the theory), but for finite samples, usually require *ad hoc* “smoothing” parameters to be specified by hand, or chosen using cross-validation. In Bayesian statistics, “smoothing” occurs naturally through the prior, but the notion of a model with no distributional assumptions at all is not particularly natural in a statistical approach that requires a well-defined distribution to be specified over all possible data generating processes.

In recent decades, however, and especially in the last decade, a field of **Bayesian nonparametrics** has grown up. Bayesian nonparametric models share with their non-Bayesian counterparts the property that the flexibility of the model increases as the sample size increases, but instead of being entirely distribution-free, Bayesian nonparametric models employ countably infinite dimensional families of distributions that can approximate large classes of distributions (e.g., all continuous distributions) arbitrarily well (Lindsay, 1995; McLachlan and Peel, 2004). The result is to combine the explicitness of assumptions in Bayesian models with the flexibility of nonparametric models.

The central innovation in this dissertation is a new nonparametric Bayesian model for sequence data, the Hierarchical Dirichlet Process Hidden Markov Model with Local Transitions (HDP-HMM-LT, or HaMMLLeT, for short). Hidden Markov Models (HMMs) are used to model observed sequence data as “noisy” realizations of an underlying sequence of discrete states. The HaMMLLeT model generalizes an existing model, the Hierarchical Dirichlet Process Hidden Markov Model (HDP-HMM), by allowing a metric structure to be specified on the state space of the underlying Markov chain, providing a mechanism to build into the state transition prior the notion that consecutive states are “similar” in some way.

A challenge for most Bayesian modeling, not least infinite-dimensional nonparametric Bayesian modeling, is developing efficient algorithms to compute (or approximate) desired functions of the posterior distribution, since for all but the simplest

models, explicit integration is intractable. For nonparametric models in particular, a key property of any practical inference algorithm is that it must be able to represent needed properties of the infinite-dimensional posterior using finite storage and processing. Thus, any time a new nonparametric model is proposed, a central challenge is introducing a tractable and efficient inference algorithm. I present a simple Gibbs sampling algorithm Geman and Geman (1984) to sample from the joint posterior of all parameters and hyperparameters of HaMMLeT as well as the latent state sequence, which is achieved by the introduction of a set of auxiliary variables with a natural interpretation under an augmentation of the stochastic process underlying HaMMLeT.

1.1.1. An overview of the dissertation

In the remainder of this chapter, I begin by discussing the general problem of model selection in the context of the bias variance tradeoff, and how Bayesian models in particular provide a natural Occam’s Razor principle to balance fidelity to data on the one hand, and simplicity on the other. Then I introduce the finite dimensional versions of the models on which HaMMLeT will build, namely mixture models and their sequential counterparts, Hidden Markov Models. In Chapter 2 I review the theory of Bayesian nonparametrics, in particular Dirichlet Processes and Hierarchical Dirichlet Processes, which underlie infinite mixture models and the infinite-state Hidden Markov Model, respectively, and then define the HDP-HMM of which HaMMLeT is a generalization. In Chapter 3 I introduce the basic theory of HaMMLeT model and an augmented stochastic process, the Markov Jump Process with Failed Transitions, which provides an interpretable set of auxiliary variables that greatly simplify posterior inference in HaMMLeT. Then in Chapters 4-6 I define some specific instantiations of HaMMLeT for a variety of kinds of data, and give experimental results comparing HaMMLeT’s inferential performance to existing models. Finally, in 7, I

give concluding thoughts and some directions for future research.

1.2. Model Selection

1.2.1. The Likelihood Function and the Conservation of Explanatory Power

In a parametric model family, $\mathcal{F} = \{f_\theta : \theta \in \Theta\}$, the principle of maximum likelihood estimation can be seen as choosing the parameter vector θ that makes the data as *unsurprising* as possible. One consequence of this principle is that models that are consistent with a broad range of possible data sets must spread their predictive probability mass over a larger volume of possible data, and thus assign a smaller probability (or probability density) to any given data set. When comparing the likelihood for a more restrictive model to a more flexible model when both are consistent with the data, then, the likelihood function will tend to prefer the restrictive one.

To see a trivial example of this idea, consider the inference problem of trying to decide what “grammatical rule” produced the following set of sequences:

lay lay dee
way way dee
fay fay dee

A few candidate rules are the following:

1. Any three syllables (call this the X Y Z rule)
2. Any repeated syllable followed by another syllable (call this X X Y)
3. Any repeated syllable followed by the syllable **dee** (X X dee)
4. Any repeated syllable rhyming with “day”, followed by the syllable **dee** (*ay *ay dee)

5. Three repetitions of the same syllable (XXX).

Which one seems to best describe the pattern? All but the XXX rule are consistent with all three examples, but if we attribute a probability model to each rule, by, for example, supposing that the content of each “free variable” is chosen uniformly from the set of viable candidates in the language, then we see that the *ay *ay dee rule is the maximum likelihood choice, being the most restrictive which is still consistent with the data.

In a sense, the likelihood principle exhibits the property of *conservation of explanatory power*: a model has a total probability mass that it can use to “place bets” on possible data sets. Models that allocate their probability mass over many data sets are less likely to be categorically wrong, but do not get as big a “reward” as a model that distributes its bet over a smaller number of possibilities. Like betting on roulette, the “vague” bet on red is reasonably likely to pay off, but results in a smaller profit than the narrow bet on 34, which is much less likely to result in a payout, but when it does, it is a large one.

This property of rewarding the most restrictive models acts as a form of “Occam’s Razor”, which is the principle that states that all else being equal we should prefer simpler explanations. But although the likelihood function rewards simpler (or at least more restrictive) explanations in these toy examples of competing models with no free parameters, it breaks down when competing models have differing numbers of “moving parts”, due to the problem of **overfitting**, that is, of confusing the stochastic idiosyncracies in a particular data set (that is, “noise”) for stable regularities in the data-generating process (that is, “signal”). Since maximum likelihood minimizes surprise, it will tend to give high scores to models that are able to shift as much variability as possible away from the stochastic noise component (which is surprising even after we have settled on a model) and the deterministic signal component (which, once we have a model, is not surprising).

1.2.2. The Bias-Variance Tradeoff

The notion of a tradeoff between model complexity and sensitivity to idiosyncracies in the data is formalized by the **bias-variance tradeoff**.

Suppose that we have a model class, \mathcal{F} , and using some data, $\{X, Y\}$, we will choose a specific model $\hat{f} \in \mathcal{F}$ that can be used to make a prediction about some future observation, y_{new} given some features, x_{new} . That is, $\hat{f}(x_{new})$ yields some predicted value \hat{y}_{new} , which we want to be as close as possible to the true but unknown value y_{new} .

Typically we cannot make perfect predictions, and there will be some error, which we might quantify using squared distance:

$$\text{Error}(x_{new}) = (\hat{f}(x_{new}) - y_{new})^2$$

Each dataset that we get yields a potentially different \hat{f} , and thus a different error, and so we can consider the *expected* error in prediction, or the **mean squared error** (MSE)

$$MSE = \mathbb{E}[(\hat{f}(x_{new}) - y_{new})^2]$$

where the expectation is taken with respect to the true (but unknown) distribution of the data.

We can decompose the MSE as follows:

$$\begin{aligned} MSE &= \mathbb{E}[(\hat{f}(x_{new}) - \mathbb{E}[\hat{f}(x_{new})] + \mathbb{E}[\hat{f}(x_{new})] - \mathbb{E}[y | x_{new}] + \mathbb{E}[y | x_{new}] - y_{new})^2] \\ &= \mathbb{V}[\hat{f}(x_{new})] + (\mathbb{E}[\hat{f}(x_{new})] - \mathbb{E}[y | x_{new}])^2 + \mathbb{V}[y_{new} | x_{new}] \\ &= \mathbb{V}[\hat{f}(x_{new})] + \text{Bias}(\hat{f}(x_{new})) + \sigma_{\varepsilon}^2(x_{new}) \end{aligned}$$


where the cross terms all cancel out. The first term in the result describes how much the prediction made by the fitted model varies depending on the specific dataset used to choose \hat{f} . The second is the squared **bias**, where the bias of a model is the average discrepancy between its prediction and the truth. Finally, σ_{ε}^2 describes the

unavoidable indeterminacy of trying to make a prediction about y_{new} using only x_{new} . This last term does not depend on the choice of model class, \mathcal{F} , but the bias and the variance do.

In general, the more complex the model class \mathcal{F} , the more room it has to fit the true pattern of the data, and so the greater its ability to be unbiased; however, this freedom often comes at the expense of greater variance: even if on average the model class yields a correct prediction, for a given dataset its prediction may be farther away from that correct “average behavior”. This phenomenon, of excessive sensitivity to idiosyncracies of the particular dataset, is known as **overfitting**. On the other hand, we can trivially produce a model with zero variance by using a constant prediction regardless of the data; but this model is presumably going to have greater bias, since unless y_{new} is not related at all to the input x_{new} , there will be values of the latter for which the corresponding y values are not centered around our constant prediction (and even if the mean of y is constant, we are unlikely to be able to intuit this without using any data).

1.2.3. Example: Polynomial Regression


We can see the tradeoff between bias and variance (and the consequences of overfitting) in the setting of fitting a polynomial curve of order d to a set of bivariate data. In Fig. 1.1, we see the resulting fit to a sample of 12 data points drawn from a polynomial model with $d = 5$ and independent Normally distributed residuals for a model class with d equal to 1, 2, and 10. The linear model has low variance, but high bias: it underfits the data. The tenth order model overfits the data, having high variance and low bias: it achieves a perfect fit on the sample used to select parameters, but we would get dramatically different predictions if the sample included the unfilled points rather than the filled points. In fact, as we see in Fig. 1.1d, the predictions made for the unfilled points by the curve fit to the filled points are poor; and the converse



(a) Linear Fit

(b) Quadratic Fit

(c) Tenth-Order Fit



(d) Errors on Training and Test Set

Figure 1.1: Polynomials of varying degrees fit using ordinary least squares (equivalently, maximum likelihood estimation) to a set of eleven points generated from a fifth order polynomial with independent Normal residuals. 1.1a–1.1c: fits of polynomial order 1, 2 and 10. Filled circles are training data, unfilled circles are data not used during fitting. 1.1d squared prediction error in and out of the training sample for polynomial orders 1 through 10.

would be true as well. The model that achieves the lowest out-of-sample prediction error is a model of intermediate complexity; in this case, just one degree away from the true model used to generate the data.

1.2.4. Balancing Fit and Complexity: Bayesian Occam’s Razor

How do we go about deciding how complex a model to use in a given setting? One common approach is to use a penalized fit statistic, in which the specific model f is chosen not to maximize raw fit to the data (e.g., by maximizing the likelihood), in which case more complex models that contain simpler models as special cases will always be chosen, but instead to maximize a *penalized* objective function that balances in-sample fit with some measure of complexity. This complexity penalty can be viewed as a means to hold down the variance of the model class; i.e., to combat overfitting.

Some examples of penalized fit measures that serve to select a subset of predictor variables to receive nonzero coefficients in linear regression are the lasso (L_1 -penalized regression), Mallows’s C_p , and various “information criteria” such as AIC and BIC (TODO: ADD CITATIONS).

Bayesian inference provides an alternative approach. Whereas the likelihood function serves as a measure of how surprising the data is given a fully specified model, Bayesian inference provides a method for quantifying how surprising the data and the specific model are together, where the data is assumed to be selected stochastically from a set of possible datasets allowed by the model (as in maximum likelihood estimation), and additionally, the *model* is treated as though it were selected stochastically from a set of possible models. At each of these choice points, if there are lots of options, any one of them becomes more surprising.

Consider the problem of choosing between two model classes, \mathcal{M}_1 and \mathcal{M}_2 , where \mathcal{M}_1 has no free parameters, and \mathcal{M}_2 has a free parameter θ , and reproduces \mathcal{M}_1 for a particular setting of $\theta = \theta_*$. The likelihood function by itself gives us no way by itself to adjudicate between the two equivalent models, \mathcal{M}_1 on the one hand, and \mathcal{M}_2

with $\theta = \theta_*$ on the other since, $p(Y | \mathcal{M}_1) = p(Y | \mathcal{M}_2, \theta_*)$. Bayes' rule yields

$$p(\mathcal{M}_1 | Y) = \frac{p(\mathcal{M}_1)p(Y | \mathcal{M}_1)}{p(Y)} \quad (1.1)$$

$$p(\mathcal{M}_2 | Y) = \frac{p(\mathcal{M}_2)p(Y | \mathcal{M}_2)}{p(Y)} \quad (1.2)$$

where in the case of \mathcal{M}_2 , we need to expand the likelihood to

$$p(Y | \mathcal{M}_2) = \int_{\Theta} p(Y | \theta, \mathcal{M}_2)p(\theta | \mathcal{M}_2)d\theta \quad (1.3)$$

which is a weighted average of specific likelihoods of the form $p(Y | \theta, \mathcal{M}_2)$ with weights given by some distribution over the parameter space. Thus, if \mathcal{M}_1 assigns a large likelihood to the observed data, even if \mathcal{M}_2 can achieve the same likelihood at θ_* , the “good value” θ_* will be diluted in the **marginal likelihood**, $p(Y | \mathcal{M}_2)$.

Hence, unlike maximum likelihood, Bayes theorem will tend to “reward” more restrictive *model classes* even in the case that the more flexible model contains the simpler one as a special case, providing a built in Occam's Razor. Of course, if there is some θ other than θ_* for which the likelihood is greater, then \mathcal{M}_2 may win out depending on whether the benefit to the likelihood outweighs the “penalty” induced by the $p(\theta | \mathcal{M}_2)$ term, which is what we want to happen, since the simplest model is only to be preferred if it does about as good a job accounting for the data as the more complex model.

I now turn to the example of clustering a set of observations into some number of groups, or, similarly, of fitting a mixture distribution to a dataset. I will introduce this problem in some detail, since it is closely related to the main topic of this dissertation, namely Hidden Markov Models, which I will introduce in the following section, and treat in more detail in Chapter 2. I introduce the clustering application here since it helps to motivate the bigger picture idea of a nonparametric model.

1.3. Clustering via a Mixture of Gaussians Model

Suppose we have a model of the form

$$f(y | \pi, \theta) = \sum_{k=1}^{\infty} \pi_k f(y | \theta_k) \quad (1.4)$$

where $\theta = (\theta_1, \theta_2, \dots)$ are some parameters governing the mixture components, and $\pi = (\pi_1, \pi_2, \dots)$ is the vector of mixing weights that sum to 1.

We could take a Bayesian approach and put priors on the vector of mixture weights π and on each θ_k .

Concretely, suppose that f is Normal, so that $\theta_k = (\mu_k, \sigma_k^2)$, and first suppose that we use a prior on π_k that fixes the number of nonzero weights to a finite value, K . Then the model is of the form

$$f(y | \pi, \theta) = \sum_{k=1}^K \pi_k \mathcal{N}(\mu_k, \sigma_k^2) \quad (1.5)$$

The π_i define a categorical distribution over K categories. The categorical family has the conjugate prior

$$f(\pi | \xi) \propto \prod_{k=1}^K \pi_k^{\xi_k - 1} \quad (1.6)$$

which is a **Dirichlet distribution**.

We can also place conjugate priors on the μ_k and σ_k , such as Normal and Inverse-Gamma, in the case of Normal components.

(Note that when y is vector valued, μ becomes a mean vector and σ^2 becomes a covariance matrix. It is also possible to define conjugate priors for these.)

Putting everything together and assuming Normal components, we have the joint prior:

$$f(\pi, \mu, \sigma^{-2}) \propto \prod_{k=1}^K \pi_k^{\xi_k - 1} e^{-\frac{1}{2\sigma_k^2}(\mu_k - \mu_0)^2} \sigma_k^{-2a_0 - 2} e^{-\sigma_k^{-2}b_0} \quad (1.7)$$

and the joint likelihood

$$f(y | \pi, \mu, \sigma^2) = \prod_{i=1}^n \sum_{k=1}^K \pi_k \sigma_k^{\frac{1}{2}} e^{-\frac{1}{2\sigma_k^2}(x_i - \mu_k)^2} \quad (1.8)$$

That sum is trouble for inference. What we can do is to represent the likelihood in two stages: instead of saying that each y_i is drawn from a weighted mixture of the K components, we can instead say that each y_i comes from precisely one component; we just do not know which one. However, we can say that y_i comes from component k with probability k . We can imagine generating data as follows:

For each $i = 1, \dots, n$

1. Draw $z_i \sim \text{Categorical}(\pi_1, \dots, \pi_K)$
2. Draw $y_i \sim f(y | \mu_{z_i}, \sigma_{z_i}^2)$

The joint likelihood for the z_i and y_i becomes

$$f(z_i, y_i | \pi, \mu, \sigma^2) = f(z_i | \pi) f(y_i | z_i, \mu, \theta) \propto \prod_{i=1}^n \pi_{z_i} (\sigma_{z_i}^2)^{-1/2} e^{-\frac{1}{2\sigma_{z_i}^2} (y_i - \mu_{z_i})^2} \quad (1.9)$$

$$= \prod_{k=1}^K \pi_k^{n_k} (\sigma_k^2)^{-\frac{n_k}{2}} e^{-\frac{1}{2\sigma_k^2} \sum_{i: z_i=k} (y_i - \mu_k)^2} \quad (1.10)$$

which yields a posterior distribution with simple conditional distributions that can be sampled from very straightforwardly using a Gibbs sampler, in which the set of posterior variables is divided into blocks, each of which is iteratively sampled from its conditional posterior given the previous state of the others, which defines a Markov Chain whose stationary distribution is the full posterior (Geman and Geman, 1984).

1.4. Clustering Sequential Data: The Hidden Markov Model

In many applications, data are not i.i.d., but have some known dependence structure. Sequential data is one such case. Suppose a dataset consists of an ordered sequence of observations, $y_t, t = 1, \dots, T$. Since unlike in the i.i.d. case we assume the order is meaningful (most typically, the index t might represent the time at which the data point was collected), it is reasonable to expect that observing the value at time t would alter the predictive distribution at time $t + 1$, even if we knew the marginal

distribution. For example, y_t might be a word or a sentence of text in a document, a snippet of a sound recording, physiological measurements collected over time. In all of these cases, we expect nearby observations either to be similar, or to be made more predictable given the last few observations.

One option would be to encode the dependencies among the data points through a conditional model of the data at time t given various combinations of data at time $t - 1, t - 2, \dots, t - s$; that is, to model the data using an order s Markov chain. However, modeling the distribution conditioned directly on observations risks an overly complex conditioning structure, since in most settings there is likely to be some idiosyncratic noise present in each observation that does not depend on time in the same way as the underlying signal. That is, quite often the distribution at t is likely to be a mixture of a temporally dependent component and a temporally independent component.

The **Hidden Markov Model**, first developed by Baum and Petrie (1966) (see also Rabiner and Juang (1986) for a classic tutorial) separates temporal dynamics of an evolving underlying (“hidden”) state from an temporally independent “noise” component by augmenting the observed time sequence $\{y_t\}, t = 1, \dots, T$ with a latent data sequence, $\{z_t\}, t = 1, \dots, T$, such that at the distribution of the observations, y_t are independent conditioned on the latent sequence, z_t , but the z_t evolve according to a Markov chain. More formally, the model is

$$z_1 \sim \text{Categorical}(\pi_{01}, \dots, \pi_{0J}) \quad (1.11)$$

$$z_t | z_{t-1} \sim \text{Categorical}(\pi_{z_{t-1}1}, \dots, \pi_{z_{t-1}J}), \quad t = 2, \dots, T \quad (1.12)$$

$$y_t | z_t \sim F(\theta_{z_t}), \quad t = 1, \dots, T \quad (1.13)$$

where the vector $\pi_0 = (\pi_{01}, \dots, \pi_{0K})$ is the initial distribution of the underlying Markov chain, the matrix $(\pi_{jj'}), j, j' = 1, \dots, J$ is the transition matrix of the Markov chain, F is a family of **emission distributions**, which is parameterized for each state in the Markov chain by the $\theta_j, j = 1, \dots, J$.

This model is nearly identical to the mixture model discussed in Sec. 1.3, the only difference being the dependence of the label, z_t on the previous label, z_{t-1} . We can indeed construct an easy to work with Bayesian HMM by placing an appropriate conjugate prior on the θ_t emission parameters (e.g., Normal-Inverse Wishart in the case where F is the family of Normal distributions), and a Dirichlet conjugate prior on π_0 and each row of π , and carry out posterior inference using Gibbs sampling: alternating between sampling the state sequence $\{z_t\}$ and the model parameters, π, π_0 and the θ_j . There is some additional difficulty introduced by the fact that the z_t indicators are no longer conditionally independent, and so we must either sample them jointly, or break up the Gibbs sampler into many additional blocks so as to sample each z_t conditioned on all the others. I will defer details of inference in this model until Chapter 2.

1.5. Model Selection in HMMs and the Mixture of Gaussians Model

It is relatively rare that we would be in a situation where we wanted to fit a mixture of Gaussians to a dataset and knew exactly how many components there were in the mixture. It is perhaps somewhat more likely that we know in advance how many states the hidden Markov chain in an HMM can visit, but there are certainly settings where we cannot or would not wish to do this. What can we do if we cannot specify K ahead of time?

One approach is to apply the logic outlined in Sec. 1.2.4 and simply put a prior distribution on K , and then calculate (or sample from) the posterior distribution over K . Then, although we can always achieve at least as large a likelihood with a larger K as with a smaller K (either by setting some mixture weights to zero, or by effectively combining multiple components into one by creating mixture components with identical means and variances), since models with larger K have more flexibility, their marginal likelihood will tend to be diluted by all of the ways they can fit the

data badly.

However, this requires doing inference separately for as many values of K as we want to consider, and for the larger values of K , this inference can be computationally expensive. It turns out that we can achieve substantially the same thing more simply by employing a single **nonparametric** model with an *infinite* number of components.

In the next chapter, I turn to the distinction between parametric and nonparametric models, develop some of the theory of **Dirichlet Process Mixture Models**, which form the basis for the infinite-state Hidden Markov Model, which I define after that, along with additional detail on Gibbs sampling in these two models.

2. BAYESIAN NONPARAMETRIC MODELS AND THE INFINITE HIDDEN MARKOV MODEL

2.1. Parametric Vs. Nonparametric Models

Models that can be identified using a finite set of values (that is, parameters) are called **parametric** models. Their complexity and expressivity is the same whether they are fit using 10 data points, 10^6 data points, or 10^{10} data points. Most canonical statistical models are parametric: the parameters in a regression model comprise the regression coefficients and the parameters of the residual distribution. The parameters of the mixture of Gaussians model comprise vector, π , of K mixing weights, and the parameters of the individual component Normal distributions: $\{\mu_k, \sigma_k^2\}, k = 1, \dots, K$.

In a parametric model, once the sample size is a few orders of magnitude larger than the number of parameters, the gains made by further increasing the sample size, which in a frequentist setting comes in the form of narrower confidence sets, and in a Bayesian setting comes in the form of reduced posterior variance, are typically negligible: the errors in prediction due to inevitable model misspecification, and practical concerns with generalizability to new data sets, overwhelm the remaining decimal places of uncertainty.

The name **nonparametric model** is something of a misnomer, in that nonparametric models do not have *no* parameters (a model with no parameters would by definition be unable to learn anything from data); rather, they have a number of parameters which grows adaptively as the sample size increases.

A common frequentist family of nonparametric models are kernel-smoothed density estimators (Rosenblatt et al., 1956; Parzen, 1962), in which the probability density of a data-generating process is estimated by taking the empirical distribution and “smoothing” it to obtain the estimated density at y by averaging together the values

nearby points using a weight kernel. As a result, the estimated density requires n values to specify it, where n is the sample size. As such, the complexity of the family of distributions in the model space grows with the sample size, unlike in a parametric model. Here, the tradeoff between bias and variance is controlled not by the complexity of the model space, but rather by the choice of smoothing bandwidth: how much of the density at each point is “borrowed” from nearby locations, with one extreme representing a bandwidth of zero, in which case the estimate is simply the empirical distribution.

It is natural to interpret the kernel density estimator as a mixture model, where there is a mixture component centered at each data point whose distribution belongs to a family defined by the kernel function: for example, the Gaussian kernel results in a mixture of Gaussians, the Epanechnikov kernel (Epanechnikov, 1969) results in a mixture of Epanechnikov distributions, and the Uniform kernel results in a mixture of Uniform distributions. On the interpretation that each component in a mixture model represents a qualitatively distinct class of data, this means that each point is viewed as qualitatively distinct, which from a Bayesian perspective is unsatisfying. Instead, we would like a model whose complexity grows more slowly than linear in the sample size, such that as we collect additional data it is always possible to encounter something qualitatively new, but where the chances of doing so diminish as we have more and more data. This can be accomplished by employing a **Dirichlet Process** as the prior on the set of mixture components.

2.2. The Dirichlet Process

The Dirichlet Process (DP) was formally defined by Ferguson (1973) as a random probability measure, G , over a \mathcal{X} equipped with the sigma-algebra \mathcal{A} , with the defining property that, for any partition of \mathcal{X} consisting of measurable sets, A_1, A_2, \dots, A_k ,

the random distribution given by the probabilities

$$\{P(A_1), P(A_2), \dots, P(A_k)\}$$

has a Dirichlet distribution.

A Dirichlet distribution is defined by a mean distribution, $\pi_1, \pi_2, \dots, \pi_k$; $\sum_i \pi_i = 1$, and a concentration parameter, α , which acts as an inverse variance parameter: as α goes to infinity, draws from the Dirichlet distribution are distributions close to the mean with increasingly high probability.

The Dirichlet Process is also defined by a mean distribution G_0 , called a **base measure**, and a concentration parameter α , but since the DP induces a Dirichlet distribution over *any* finite partition of \mathcal{X} , the mean distribution must be defined on all measureable sets in \mathcal{A} . We will write

$$G \sim \text{DP}(\alpha G_0) \tag{2.1}$$

to indicate that the random measure G is distributed according to a Dirichlet Process with base measure G_0 and concentration parameter α . Then, concretely, we have

$$P(A_1), P(A_2), \dots, P(A_k) \sim \text{Dirichlet}(\alpha G_0(A_1), \dots, \alpha G_0(A_k))$$

An important probability of the DP is that the resulting measure is discrete almost surely, and hence it is a sensible choice as a prior on mixture components. An equally important property when it comes to using a DP as a Bayesian prior is that it is a *conjugate prior* to a discrete likelihood. That is, if n observations, Y_1, \dots, Y_n , are drawn from some unknown discrete probability measure G , and the prior employed for G is a Dirichlet Process, then the posterior distribution on G is also a DP, whose base measure is a weighted sum of the prior base measure, G_0 , and the empirical distribution, \hat{F}_n :

$$G | Y \sim \alpha G_0 + n \hat{F}_n \tag{2.2}$$

and whose concentration parameter is simply $\alpha + n$.

2.2.1. The Normalized Gamma Process representation of the DP

Ferguson (1973) also showed that the Dirichlet Process arises by normalizing a **Gamma Process**. A Gamma Process is a stochastic point process on $\mathbb{R}^+ \times \Theta$ which is defined by a **Lévy intensity measure**:

$$\nu(d\pi, d\theta) = \alpha \pi^{-1} e^{-\pi} d\pi G_0(d\theta) \quad (2.3)$$

A realization of a Gamma process is a collection of point masses $\{\pi_k, \theta_k\}$, where $\pi_k \in \mathbb{R}^+$ is the mass associated with the point at location $\theta_k \in \Theta$. This collection can be used to define a measure, μ , on Θ , where

$$\mu(A) = \sum_{k: \theta_k \in A} \pi_k \quad (2.4)$$

The number $n(A)$ of point masses in a region $A \subset \mathbb{R}^+ \times \Theta$ is distributed as

$$n(A) \sim \text{Poisson}\left(\int_A \nu(d\pi, d\theta)\right) \quad (2.5)$$

The Lévy intensity measure of the Gamma process satisfies conditions to guarantee that the sum of all of the π_k weights is finite with probability 1, and therefore the measure G defined by

$$G = \frac{\mu}{\sum_k \pi_k} \quad (2.6)$$

is a valid probability measure. Ferguson (1973) showed that this probability measure is a DP with base measure G_0 and concentration α , where these are the measure and parameter used in defining the Lévy intensity of the Gamma process.

Although the formal definition of the DP is well-defined, and although the Normalized Gamma Process representation guarantees existence of the DP, neither of these is terribly useful in *constructing* a DP, which limits the usefulness of the DP in applied modeling. Fortunately, a constructive definition of the DP was discovered by Sethuraman (1994), using what is known as a **Stick-Breaking Process**.

2.2.2. The Stick-Breaking Process construction of the DP

As shown by Sethuraman (1994), we can generate a draw from a Dirichlet Process by iteratively sampling the π_k weights, and then placing a point mass with weight π_k at a location in Θ independently drawn from G_0 . This process is called a **Stick-Breaking Process**. By using the following algorithm to select the stick weights, the resulting collection of point masses has a Dirichlet Process.

Having drawn $k - 1$ point masses $(\pi_1, \theta_1), \dots, (\pi_{k-1}, \theta_{k-1})$,

1. Draw $\tilde{\pi}_k \sim \text{Beta}(1, \alpha)$.
2. Set $\pi_k = \tilde{\pi}_k \prod_{k'=1}^{K-1} \tilde{\pi}_{k'}$.
3. Draw $\theta_k \sim G_0$.

where, when $k = 1$, the null product in step 2 is 1.

The choice of G_0 and α determine the resulting DP.

2.2.3. The Chinese Restaurant Process

A useful construction for doing inference in a Dirichlet Process-based model is based on the metaphor of customers sharing food at a Chinese restaurant, which is normally described as follows.

One by one, customers enter a Chinese restaurant and either sit at an unoccupied table and order a dish for the table, or join a table with other customers and share whatever dish is at the table. The first customer necessarily starts a new table. Subsequent customers join table k with probability proportional to n_k , where n_k is the number of customers currently seated at that table, and sit at an unoccupied table with probability proportional to a parameter α .

More formally, the $n + 1$ th customer to enter the restaurant is assigned to table k

according to the distribution

$$P(t_{n+1} = k) = \begin{cases} \frac{n_k}{n+\alpha} & k = 1, \dots, K \\ \frac{\alpha}{n+\alpha} & k = K + 1 \end{cases} \quad (2.7)$$

For each new table, a dish, θ_k is sampled from a base distribution, G_0 .

It turns out that the distribution of the collection of assignments of dishes to customers is the same as the marginal distribution of draws from a random measure G which is distributed $\text{DP}(\alpha G_0)$ (see Teh (2011) for a proof of this fact as well as a review of the theory of DPs in general).

Notice that the Chinese Restaurant Process (CRP) has the property that the more often a dish has already been selected, the more likely it is to be selected again: that is, it has a “rich get richer” quality. This makes sense in terms of the marginal distribution of draws from a DP-distributed random measure since, the more observations there are at a particular location, the stronger the evidence that the mass under G at that location is large, and hence, the larger the posterior predictive probability at that location (marginalizing over G).

2.2.4. The Dirichlet Process Mixture Model

We are now ready to define a nonparametric version of a Bayesian Gaussian Mixture Model which replaces the Dirichlet distribution prior on the collection of mixture components from the fixed K mixture model with a Dirichlet Process prior.

Reiterating the model initially defined in (1.4), suppose we have a model of the form

$$f(y | \pi, \theta) = \sum_{k=1}^{\infty} \pi_k f(y | \theta_k) \quad (2.8)$$

where $\theta = (\theta_1, \theta_2, \dots)$ are some parameters governing the mixture components, and $\pi = (\pi_1, \pi_2, \dots)$ is the vector of mixing weights that sum to 1. We now drop the assumption that all but finitely many of the π_k are zero, and instead adopt as a prior:

$$\{(\pi_k, \theta_k)_{k=1}^{\infty}\} \sim \text{DP}(\alpha G_0) \quad (2.9)$$

As before, we introduce indicator variables, $\{z_i\}_{i=1}^n$ so that we can write

$$P(z_i = k) = \pi_k, \quad i = 1, \dots, n; k = 1, 2, \dots \quad (2.10)$$

$$y_i | z_i \sim F(\theta_{z_i}) \quad (2.11)$$

where F is a parametric family parameterized by θ . For example, if F is Normal, then θ might represent the mean vector and covariance matrix, and we might choose G_0 to be a Normal-Inverse Wishart distribution (or a non-conjugate choice as the application suggests).

2.2.5. A Gibbs Sampler for DP Mixture Models

Given data, $y = (y_1, \dots, y_n)$ modeled using the mixture model in (2.8) with the prior in (2.9), we want to be able to make inferences about the parameters, (π_k, θ_k) . As with all but the simplest models, the full posterior is not amenable to exact analysis, and so we must resort to approximation to compute quantities of interest. There are two dominant approximation methods in Bayesian inference. One is variational inference, in which the true posterior is replaced by a distribution which is more amenable to analysis and which is in some sense as close as possible to the true distribution (usually the objective is to minimize the KL divergence from the target to the approximation). The second popular approximation technique is Markov Chain Monte Carlo (MCMC), in which integrals involving the posterior are computed based on a sample from the posterior drawn using a Markov Chain whose transition kernel is chosen so as to yield the true posterior as the unique stationary distribution. In a nutshell, variational Bayes provides an exact calculations based on an approximate distribution, whereas MCMC provides approximate calculations based on the true posterior. I will focus on MCMC in this dissertation Fox and Roberts (2012) for a tutorial on variational Bayes generally and Blei and Jordan (2006) and Kurihara et al. (2007) for work on variational methods for Dirichlet Process mixture models in particular.

As with the finite mixture model, MCMC inference is greatly simplified by sampling over the z_i indicators as well as over the mixture component parameters themselves (indeed, in some applications, these indicators may be the variables of primary interest).

2.3. An Infinite State HMM

2.3.1. The Hierarchical Dirichlet Process

2.3.2. The HDP-HMM

2.4. Inference in Finite and Infinite State HMMs

2.4.1. The Forward Backward Algorithm

2.4.2. Two Gibbs Samplers for the HDP-HMM

2.4.3. Beam Sampling in the HDP-HMM

3. HAMMLET: AN INFINITE HIDDEN MARKOV MODEL WITH LOCAL TRANSITIONS

In this chapter, I describe a generalization of the Hierarchical Dirichlet Process Hidden Markov Model Teh et al. (2006) discussed in Chapter 2 which introduces a notion of latent similarity between pairs of hidden states, such that transitions are a priori more likely to occur between states that are deemed “similar”. This is achieved by placing a kernel function on the space of state parameters which returns for each pair of states a value between 0 and 1, representing the degree to which they are similar (in whatever sense is desired for the application at hand), and scaling HDP-generated transition probabilities by the similarity between states. I will refer to this model as the Hierarchical Dirichlet Process Hidden Markov Model with Local Transitions (HDP-HMM-LT, or HaMMLeT). Although this achieves the goal of selectively increasing the probability of transitions between similar states, inference is made more complicated since, unlike in the “vanilla” HDP-HMM, the posterior measure over transition distributions is no longer a Dirichlet Process — that is, the prior is no longer conjugate to the state sequence likelihood.

I will present an alternative representation of this process that facilitates inference with an auxiliary variable scheme with the following interpretation: The discrete time chain is recast as a continuous time Markov Process in which: (1) some jump attempts fail, (2) the probability of success is proportional to the similarity between the source and destination states, (3) only successful jumps are observed, and (4) the time elapsed between jumps, as well as the number of unsuccessful jump attempts, are latent variables that are sampled during MCMC inference. By introducing these auxiliary latent variables, nearly all conditional distributions in the model are members of an exponential family, admitting exact Gibbs sampling. The only exception is the set of similarities, which are defined in an application-specific way, and require

application-specific inference methods. I present results for a few different choices in the experiments in Chs 4 through 6.

The motivating domain for this model is natural language text, in which sentences in a document are associated with a set of “topics”, the topic sets in successive sentences have a high degree of overlap, even when they are not identical (so that neighboring topic vectors are similar), and in which there may be a high degree of correlation between topics, so that modeling the entry and disappearance of each topic independently is undesirable. The topic vectors are represented using binary vectors, indicating which topics are “active” in the sentence, and to constrain the dynamics governing latent state transitions so that transitions between similar topic vectors are *a priori* more likely, but where certain topics tend to occur together. The latter property makes an ordinary factorial HMM undesirable.

In the remainder of this chapter, I first review the transition dynamics in the “vanilla” HDP-HMM so that it is easier to see how the proposed HDP-HMM-LT differs; I then define the generative process of the HDP-HMM-LT; finally, I introduce the “Markov Jump Process With Failed Transitions” augmented data representation, which gives rise to a natural Gibbs sampler for the HDP-HMM-LT model.

3.1. Transition Dynamics in the HDP-HMM

The conventional HDP-HMM (Teh et al., 2006) is based on a Hierarchical Dirichlet Process which is defined as follows:

Each of a countably infinite set of states, indexed by j , receives a parameter vector $\theta_j \in \Theta$, according to base measure H . A top-level weight distribution, β , is drawn from a stick-breaking process with parameter $\gamma > 0$, so that state j has overall weight

β_j , and an emission distribution which is parameterized by θ_j .

$$\theta_j \stackrel{i.i.d.}{\sim} H \quad (3.1)$$

$$\beta \sim \text{GEM}(\gamma) \quad (3.2)$$

The actual transition distribution from state j , denoted by π_j is then drawn from a Dirichlet Process with concentration parameter α and base measure β :

$$\pi_j \stackrel{i.i.d.}{\sim} \text{DP}(\alpha\beta) \quad j = 1, 2, \dots \quad (3.3)$$

The hidden state sequence is then generated according to the π_j . Let z_t be the index of the chain's state at time t . Then we have

$$z_t \mid z_{t-1}, \pi_{z_{t-1}} \sim \pi_{z_{t-1}} \quad t = 1, 2, \dots, T \quad (3.4)$$

where T is the length of the data sequence.

Finally, the emission distribution for state j is a function of θ_j , so that we have

$$y_t \mid z_t, \theta_{z_t} \sim F(\theta_{z_t}) \quad (3.5)$$

A shortcoming of this model is that the generative process does not take into account the fact that the set of source states is the same as the set of destination states: that is, that the distribution π_j has an element which corresponds to state j . Put another way, there is no special treatment of the diagonal of the transition matrix, so that self-transitions are no more likely *a priori* than transitions to any other state.

The Sticky HDP-HMM (Fox et al., 2008) addresses this issue by adding an extra mass of κ at location j to the base measure of the DP that generates π_j . That is, they replace (4.2) with

$$\pi_j \sim \text{DP}(\alpha\beta + \kappa\delta_j). \quad (3.6)$$

An alternative model is presented by Johnson and Willsky (2013), wherein state duration distributions are modeled separately, and ordinary self-transitions are ruled

out. In both of these models, auxiliary latent variables are introduced to simplify conditional posterior distributions and facilitate Gibbs sampling. However, while both of these models have the useful property that self-transitions are treated as “special”, they contain no notion of similarity for pairs of states that are not identical: in both cases, when $j \neq j'$, the prior probability of transitioning from j to j' depends only on the top-level stick weight associated with state j' , and not on the identity or parameters of the previous state j .

3.2. An HDP-HMM With Local Transitions

The goal of the proposed model is to add to the transition model the concept of a transition to a “nearby” state, where nearness of j and j' may be defined deterministically or stochastically in terms of the emission parameters, θ_j and $\theta_{j'}$, or may be based on a latent state geometry that is a priori independent of the emission distribution.

In order to accomplish this, we first consider an alternative construction of the transition distributions, based on the Normalized Gamma Process representation of the Dirichlet Process (Ferguson, 1973).

Notational Conventions In the definitions and derivations that follow, I will adopt the convention that variables written with no subscript, such as θ , β , and π , represent the collection of all corresponding subscripted values: for example, θ is the vector $(\theta_1, \theta_2, \dots)$, and π is the matrix $(\pi_{jj'})$. For variables that represent counts, I will use a dot in the subscript to represent a sum over corresponding individual counts; for example, $n_{j\cdot}$ is used to represent $\sum_{j'} n_{jj'}$, and $n_{\cdot\cdot}$ means $\sum_j \sum_{j'} n_{jj'}$.

3.2.1. A Normalized Gamma Process representation of the HDP-HMM

Define a random measure, $\mu = \sum_{j=1}^{\infty} \pi_j \delta_{\theta_j}$, where

$$\pi_j \stackrel{ind}{\sim} \text{Gamma}(\omega_j, 1) \quad (3.7)$$

$$T = \sum_{j=1}^{\infty} \pi_j \quad (3.8)$$

$$\tilde{\pi}_j = \frac{\pi_j}{T} \quad (3.9)$$

$$\theta_j \stackrel{i.i.d}{\sim} H \quad (3.10)$$

and subject to the constraint that $\sum_{j \geq 1} \omega_j < \infty$, which ensures that $T < \infty$ almost surely, since

$$T \sim \text{Gamma}\left(\sum_j \omega_j, 1\right).$$

As shown by Paisley et al. (2011), for fixed $\{\omega_j\}$ and $\{\theta_j\}$, μ is distributed as a Dirichlet Process with base measure $\nu = \sum_{j=1}^{\infty} \omega_j \delta_{\theta_j}$. If we draw β from a stick-breaking process and then draw a series $\{\mu_m\}_{m=1}^M$ of i.i.d. random measures from the above process, setting $w = \alpha\beta$ for some $\alpha > 0$, then this defines a Hierarchical Dirichlet Process. If, moreover, there is one μ_m associated with every state j , then we obtain the HDP-HMM.

We can thus write

$$\beta \sim \text{GEM}(\gamma) \quad (3.11)$$

$$\theta_j \stackrel{i.i.d.}{\sim} H \quad (3.12)$$

$$\pi_{jj'} \stackrel{ind}{\sim} \text{Gamma}(\alpha\beta_{j'}, 1) \quad (3.13)$$

$$T_j = \sum_{j'=1}^{\infty} \pi_{jj'} \quad (3.14)$$

$$\tilde{\pi}_{jj'} = \frac{\pi_{jj'}}{T_j}, \quad (3.15)$$

where γ and α are prior concentration hyperparameters for the two DP levels,

$$p(z_t | z_{t-1}, \pi) = \tilde{\pi}_{z_{t-1} z_t} \quad (3.16)$$

and the observed data $\{y_t\}_{t \geq 1}$ is distributed as

$$y_t \mid z_t \stackrel{\text{ind}}{\sim} F(\theta_{z_t}) \quad (3.17)$$

for some family, F of probability measures indexed by values of θ .

3.2.2. Promoting “Local” Transitions

In the preceding formulation, the transition distributions $\{\pi_j\}_{j=1}^\infty$ are independent conditioned on the top-level weights, β . Our goal is to relax this assumption, in order to allow for the possibility that there may be correlations between these distributions. We achieve this by introducing a notion of a geometry on the state space; that is, that some pairs of states are “near” to each other, and will thus tend to have similar transition probabilities associated with them.

We associate with each state a location, $\ell_j \in \mathcal{L}$, and define a “similarity function”, $\phi : \mathcal{L} \times \mathcal{L} \rightarrow [0, 1]$, which returns for any pair of locations a measure of how “close” they are.

In order to remain agnostic about the extent to which ϕ is based on the emission parameters, I will use θ_j to represent the emission parameters for state j , and assume that ℓ_j determines θ_j , but that ϕ may be based on any part of ℓ_j , which may include some, all, or none of the information contained in θ_j .

I will also use the shorthand $\phi_{jj'}$ to represent $\phi(\ell_j, \ell_{j'})$, for the sake of readability; but the reader should keep in mind that whenever ϕ appears in a conditional distribution, it is a constant if and only if ℓ is being conditioned on.

We can generalize the generative process defined in (3.11)-(3.15) as follows:

$$\beta \sim \mathbf{GEM}(\gamma) \quad (3.18)$$

$$\ell_j \stackrel{i.i.d}{\sim} H \quad (3.19)$$

$$\pi_{jj'} \mid \beta, \theta \sim \mathbf{Gamma}(\alpha\beta_{j'}, \phi_{jj'}^{-1}) \quad (3.20)$$

$$T_j = \sum_{j'=1}^{\infty} \pi_{jj'} \quad (3.21)$$

$$\tilde{\pi}_{jj'} = \frac{\pi_{jj'}}{T_j} \quad (3.22)$$

$$y_t \mid z_t \stackrel{ind}{\sim} F(\theta_{z_t}) \quad (3.23)$$

where H is now a measure on \mathcal{L} . In this new formulation, the expected value of $\pi_{jj'}$ is $\alpha\beta_{j'}\phi_{jj'}$. Since a similarity between one object and another should not exceed the similarity between an object and itself, we will assume that $\phi_{jj'} = 1$ if $j = j'$. We will also assume that the similarity function is symmetric, so that $\phi_{jj'} = \phi_{j'j}$ for all j, j' . As we will see, either or both of these assumptions could be relaxed if desired in a particular application, but derivations presented here will make both.

The above model is equivalent to simply drawing the $\pi_{jj'}$ as in (3.11) and scaling each one by $\phi_{jj'}$ prior to normalization, since it is a general property of Gamma distributions that, if $X \sim \mathbf{Gamma}(a, b)$, then $cX \sim \mathbf{Gamma}(a, b/c)$.

Unfortunately, this formulation complicates inference significantly, compared to the ordinary HDP-HMM, as the introduction of non-constant rate parameters to the prior on π destroys the conjugacy between π and z , and worse, the conditional likelihood function for π contains a sum which renders all entries within a row mutually dependent *a posteriori*.

3.3. The HDP-HMM-LT as a continuous-time Markov Jump Process with “failed” jumps

We can gain stronger intuition, as well as simplify posterior inference, by re-casting the HDP-HMM-LT described in the last section as a continuous time Markov Jump Process where some of the attempts to jump from one state to another fail, and where the failure probability increases as a function of the “distance” between the states.

Let ϕ be defined as in the last section, and let β , θ and π be defined as in the Normalized Gamma Process representation of the ordinary HDP-HMM. Specifically,

$$\beta \sim \text{GEM}(\gamma) \quad (3.24)$$

$$\ell_j \stackrel{i.i.d}{\sim} H \quad (3.25)$$

$$\pi_{jj'} \mid \beta \sim \text{Gamma}(\alpha\beta_{j'}, 1) \quad (3.26)$$

Now suppose that when the process is in state j , jumps to state j' are made at rate $\pi_{jj'}$. This defines a continuous-time Markov Process where the off-diagonal elements of the transition rate matrix are the off diagonal elements of the π matrix. In addition, self-jumps are allowed, and occur with rate π_{jj} .

If we only observe the jumps and not the durations between jumps, this is an ordinary Markov chain, whose transition matrix is obtained by appropriately normalizing the rows of π . If we do not observe the jumps themselves, but instead an observation is generated once per jump from a distribution that depends on the state being jumped to, then we have an ordinary HMM.

We modify this process as follows. Suppose that each jump attempt from state j to state j' has a chance of failing, which is an increasing function of the “distance” between the states. In particular, let the success probability be $\phi_{jj'}$ (recall that we assumed above that $0 \leq \phi_{jj'} \leq 1$ for all j, j'). Then, the rate of successful jumps from j to j' is $\pi_{jj'}\phi_{jj'}$, and the corresponding rate of unsuccessful jump attempts is $\pi_{jj'}(1 - \phi_{jj'})$. To see this, denote by $N_{jj'}$ the total number of jump attempts to j' in a

unit interval of time spent in state j . Since we are assuming the process is Markovian, the total number of attempts is $\text{Poisson}(\pi_{jj'})$ distributed. Conditioned on $N_{jj'}$, $n_{jj'}$ will be successful, where

$$n_{jj'} \mid N_{jj'} \sim \text{Binom}(N_{jj'}, \phi_{jj'}) \quad (3.27)$$

It is easy to show (and well known) that the marginal distribution of $n_{jj'}$ is $\text{Poisson}(\pi_{jj'}\phi_{jj'})$, and the marginal distribution of $N_{jj'} - n_{jj'}$ is $\text{Poisson}(\pi_{jj'}(1 - \phi_{jj'}))$. The rate of successful jumps from state j overall is then $T_j := \sum_{j'} \pi_{jj'}\phi_{jj'}$.

Let t index jumps, so that z_t indicates the t th state visited by the process (counting self-jumps as marking a transition to a new time step). Given that the process is in state j at discretized time t (that is, $z_t = j$), it is a standard property of Markov Processes that the probability that the destination z_{t+1} of the first successful jump is independent of the time since the last jump, and the probability that $z_{t+1} = j'$ is proportional to the rate, $\pi_{jj'}\phi_{jj'}$.

Let τ_t indicate the time elapsed between the $t - 1$ th and t th successful jump (where we assume that the first observation occurs when the first successful jump from a “dummy” initial state is made). We have

$$\tau_t \mid z_{t-1} \sim \text{Exponential}(T_{z_{t-1}}) \quad (3.28)$$

where τ_t is independent of z_t .

During this period, there will be some number of unsuccessful attempts to jump to each state, where the rate of unsuccessful jump attempts from to state j' is given by $\pi_{z_{t-1}}(1 - \phi_{z_{t-1}j'})$. Denote by $\tilde{q}_{j't}$ the number of unsuccessful jump attempts to j' during the interval with duration τ_t . Then we have

$$\tilde{q}_{j't} \mid z_{t-1}, \tau_t \sim \text{Poisson}(\tau_t \pi_{z_{t-1}j'}(1 - \phi_{z_{t-1}j'})) \quad (3.29)$$

Define the following additional variables

$$\mathcal{T}_j = \{t \mid z_{t-1} = j\} \quad (3.30)$$

$$q_{jj'} = \sum_{t \in \mathcal{T}_j} \tilde{q}_{j't} \quad (3.31)$$

$$u_j = \sum_{t \in \mathcal{T}_j} \tau_t. \quad (3.32)$$

In addition, let $Q = (q_{jj'})_{j,j' \geq 1}$ be the matrix of unsuccessful jump attempt counts, and $u = (u_j)_{j \geq 1}$ be the vector whose j th entry is the total time spent in state j .

Since each of the τ_t with $t \in \mathcal{T}_j$ are i.i.d. $\text{Exponential}(T_j)$, and since the sum of n Exponential random variables with shared scale λ has a $\text{Gamma}(n, \lambda)$ distribution, we have

$$u_j \mid z, \pi, \ell \stackrel{\text{ind}}{\sim} \text{Gamma}(n_{j\cdot}, T_j) \quad (3.33)$$

where we define $n_{j\cdot} = \sum_{j'} n_{jj'}$, where $n_{jj'}$ is the number of successful jumps from state j to j' and $n_{j\cdot}$ is the total number of times that state j is visited.

Moreover, since the $\tilde{q}_{j't}$ with $t \in \mathcal{T}_j$ are Poisson distributed, the total number of failed attempts in the total duration u_j is

$$q_{jj'} \stackrel{\text{ind}}{\sim} \text{Poisson}(u_j \pi_{jj'} (1 - \phi_{jj'})). \quad (3.34)$$

Thus if we marginalize out the individual τ_t and $\tilde{q}_{j't}$, we have a joint distribution over z , u , and Q , conditioned on the transition rate matrix π and the success

probability matrix ϕ , which is

$$p(z, u, Q | \pi, \ell) = \left(\prod_{t=1}^T p(z_t | z_{t-1}) \right) \prod_j p(u_j | z, \pi, \ell) \prod_{j'} p(q_{jj'} | u_j \pi_{jj'}, \phi_{jj'}) \quad (3.35)$$

$$= \left(\prod_t \frac{\pi_{z_{t-1}z_t} \phi_{z_{t-1}z_t}}{T_{z_{t-1}}} \right) \prod_j \frac{T_j^{n_{j\cdot}}}{\Gamma(n_{j\cdot})} u_j^{n_{j\cdot}-1} e^{-T_j u_j} \quad (3.36)$$

$$\times \prod_{j'} e^{-u_j \pi_{jj'} (1-\phi_{jj'})} u_j^{q_{jj'}} \pi_{jj'}^{q_{jj'}} (1-\phi_{jj'})^{q_{jj'}} (q_{jj'}!)^{-1} \quad (3.37)$$

$$= \prod_j \Gamma(n_{j\cdot})^{-1} u_j^{n_{j\cdot}+q_{j\cdot}-1} \quad (3.38)$$

$$\times \prod_{j'} \pi_{jj'}^{n_{jj'}+q_{jj'}} \phi_{jj'}^{n_{jj'}} (1-\phi_{jj'})^{q_{jj'}} e^{-\pi_{jj'} \phi_{jj'} u_j} e^{-\pi_{jj'} (1-\phi_{jj'}) u_j} (q_{jj'}!)^{-1} \quad (3.39)$$

$$= \prod_j \Gamma(n_{j\cdot})^{-1} u_j^{n_{j\cdot}+q_{j\cdot}-1} \prod_{j'} \pi_{jj'}^{n_{jj'}+q_{jj'}} \phi_{jj'}^{n_{jj'}} (1-\phi_{jj'})^{q_{jj'}} e^{-\pi_{jj'} u_j} (q_{jj'}!)^{-1} \quad (3.40)$$

3.3.1. An HDP-HSMM-LT modification

In any Hidden Markov Model, the distribution of the number of time steps for which a given hidden state persists is by definition a Geometric distribution, where the “failure” parameter is the relevant entry on the diagonal of the transition matrix. The HDP-HMM and the HDP-HMM-LT as defined above are no exception. Although the Sticky HDP-HMM Fox et al. (2008) and the LT generalization presented here provide mechanisms for which the diagonal entries of the transition matrix will tend to have greater mass than the off-diagonal entries, they do not alter the Markovian assumption, which implies Geometric durations.

The HDP Hidden Semi-Markov Model (HDP-HSMM; Johnson and Willsky (2013)) gets around this restriction directly, by treating self-transitions as fundamentally distinct from all other transitions, and modeling state persistence durations directly. Should it be desirable to combine this property with the general similarity bias of the HDP-HMM-LT, it is trivial to modify the LT model to incorporate a separate

duration model. We can simply fix the diagonal elements of π to be zero, and allow D_t observations to be emitted *i.i.d.* $F(\theta_{z_t})$ at jump t , where

$$D_t | z \stackrel{ind}{\sim} g(\omega_{z_t}) \quad \omega_j \stackrel{i.i.d}{\sim} G \quad (3.41)$$

The likelihood then includes the additional term for the D_t , and the only inference step which is affected is that, instead of sampling z alone, we sample z and the D_t jointly, by defining

$$z_s^* = z_{\max\{T | s \leq \sum_{t=1}^T D_t\}} \quad (3.42)$$

where s ranges over the total number of observations, and associating a y_s observation sequence with each z_s^* . Inferences about ϕ are not affected, since the diagonal elements are assumed to be 1 anyway.

In the HDP-HSMM as it is presented in Johnson and Willsky (2013), zeroing out the diagonal of the transition matrix to isolate self-transitions to the separate duration model necessitates renormalization of the other entries so that the rows of the transition matrix are probability distributions. As a result, the conditional posterior for the matrix is no longer in an exponential family. To deal with this, Johnson and Willsky introduce auxiliary variables which can be interpreted as the diagonal entries of the transition matrix prior to zeroing out, as well as the number of self-transitions that would have occurred for each state had the transition matrix diagonal governed self-transitions instead of the separate duration model. Conditioned on these auxiliary variables, conjugacy between the transition matrix prior and likelihood is restored, and Gibbs sampling is able to proceed with exponential family updates.

In the LT model, on the other hand, we are already representing the transition matrix in unnormalized form, having rendered the entries conditionally independent given the z state sequence and the u holding times, so we are free to clamp the diagonal entries at zero without introducing a need for additional auxiliary variables in order to achieve semi-Markov dynamics.

3.3.2. Summary

I have defined the following augmented generative model for the HDP-H(S)MM-LT:

$$\beta \sim \text{GEM}(\gamma) \quad (3.43)$$

$$\ell_j \stackrel{i.i.d}{\sim} H \quad (3.44)$$

$$\pi_{jj'} | \beta \sim \text{Gamma}(\alpha\beta_{j'}, 1) \quad (3.45)$$

$$z_t | z_{t-1}, \pi, \ell \sim \sum_j \left(\frac{\pi_{z_{t-1}j} \phi_{z_{t-1}j}}{\sum_{j'} \pi_{z_{t-1}j'} \phi_{z_{t-1}j'}} \right) \delta_j \quad (3.46)$$

$$u_j | z, \pi, \ell \stackrel{ind}{\sim} \text{Gamma}(n_j, \sum_{j'} \pi_{jj'} \phi_{jj'}) \quad (3.47)$$

$$q_{jj'} | u, \pi, \ell \stackrel{ind}{\sim} \text{Poisson}(u_j(1 - \phi_{jj'})\pi_{jj'}) \quad (3.48)$$

$$y_t | z, \ell \sim F(\theta_{z_t}) \quad (3.49)$$

If we are using the HSMM variant, then we simply fix π_{jj} to 0 for each j , draw

$$\omega_j \stackrel{i.i.d}{\sim} G \quad (3.50)$$

$$D_t | z \stackrel{ind}{\sim} g(\omega_{z_t}), \quad (3.51)$$

for chosen G and g , set

$$z_s^* = z_{\max\{T | s \leq \sum_{t=1}^T D_t\}} \quad (3.52)$$

and replace (3.49) with

$$\mathbf{y}_s | z, \theta \sim F(\theta_{z_s^*}) \quad (3.53)$$

3.4. MCMC Inference in the “Failed Jumps” Representation

I develop a Gibbs sampling algorithm based on the Markov Process with Failed Jumps representation, augmenting the data with the duration variables u , the failed jump attempt count matrix, Q , as well as additional auxiliary variables which we will define below. In this representation the transition matrix is not modeled directly, but is a

function of the unscaled transition matrix π and the similarity matrix ϕ . The full set of variables is partitioned into three blocks: (1) $\{\gamma, \alpha, \beta, \pi\}$, (2) $\{z, u, Q, \xi\}$, and (3) $\{\ell\}$, where ξ is a placeholder for an additional set of auxiliary variables that will be introduced below. The variables in each block are sampled jointly conditioned on the other two blocks. In some of the applications described in later chapters, blocks (2) and (3) can be sampled jointly, and if desired, the parameters of the ϕ function can be given priors and sampled as a separate block.

Since we are representing the transition matrix of the Markov chain explicitly, we approximate the stick-breaking process that produces β using a finite Dirichlet distribution with a number of components larger than we expect to need, forcing the remaining components to have zero weight.

Let J indicate the maximum number of states. Then, we approximate (3.24) with

$$\beta \mid \gamma \sim \text{Dirichlet}(\gamma/J, \dots, \gamma/J) \quad (3.54)$$

This distribution converges weakly to the Stick-Breaking Process as $J \rightarrow \infty$ Ishwaran and Zarepour (2000). In practice, J is large enough when the vast majority of the probability mass in β is allocated to a strict subset of components, or when the latent state sequence z never uses all J available states, indicating that the data is well described by a number of states less than J .

3.4.1. Sampling π , β , α and γ

The joint conditional over γ , α , β and π given z , u , Q , ξ and θ will factor as

$$p(\gamma, \alpha, \beta, \pi \mid z, u, Q, \xi, \theta) = p(\gamma \mid \xi)p(\alpha \mid \xi)p(\beta \mid \gamma, \xi)p(\pi \mid \alpha, \beta, \theta, z) \quad (3.55)$$

I will derive these four factors in reverse order.

Sampling π The entries in π are conditionally independent given α and β , so we have the prior

$$p(\pi \mid \beta, \alpha) = \prod_j \prod_{j'} \Gamma(\alpha\beta_{jj'})^{-1} \pi_{jj'}^{\alpha\beta_{jj'}-1} \exp(-\pi_{jj'}), \quad (3.56)$$

and the likelihood given augmented data $\{z, u, Q\}$ given by (3.40). Combining these, we have

$$p(\pi, z, u, Q \mid \beta, \alpha, \theta) = \prod_j u_j^{n_j + q_j - 1} \prod_{j'} \Gamma(\alpha\beta_{jj'})^{-1} \pi_{jj'}^{\alpha\beta_{jj'} + n_{jj'} + q_{jj'} - 1} \quad (3.57)$$

$$\times e^{-(1+u_j)\pi_{jj'}} \phi_{jj'}^{n_{jj'}} (1 - \phi_{jj'})^{q_{jj'}} (q_{jj'}!)^{-1} \quad (3.58)$$

Conditioning on everything except π , we get

$$p(\pi \mid Q, u, z, \beta, \alpha, \theta) \propto \prod_j \prod_{j'} \pi_{jj'}^{\alpha\beta_{jj'} + n_{jj'} + q_{jj'} - 1} \exp(-(1 + u_j)\pi_{jj'}) \quad (3.59)$$

and thus we see that the $\pi_{jj'}$ are conditionally independent given u, z and Q , and distributed according to

$$\pi_{jj'} \mid n_{jj'}, q_{jj'}, \beta_{jj'}, \alpha \stackrel{ind}{\sim} \text{Gamma}(\alpha\beta_{jj'} + n_{jj'} + q_{jj'}, 1 + u_j) \quad (3.60)$$

Sampling β Consider the conditional distribution of β having integrated out π . The prior density of β from (3.54) is

$$p(\beta \mid \gamma) = \frac{\Gamma(\gamma)}{\Gamma(\frac{\gamma}{J})^J} \prod_j \beta_j^{\frac{\gamma}{J} - 1} \quad (3.61)$$

After integrating out π in (3.57), we have

$$p(z, u, Q | \beta, \alpha, \gamma, \theta) = \prod_{j=1}^J u_j^{-1} \prod_{j'=1}^J u^{n_{jj'}+q_{jj'}-1} (1+u_j)^{-(\alpha\beta_{j'}+n_{jj'}+q_{jj'})} \quad (3.62)$$

$$\times \frac{\Gamma(\alpha\beta_{j'} + n_{jj'} + q_{jj'})}{\Gamma(\alpha\beta_{j'})} \phi_{jj'}^{n_{jj'}} (1 - \phi_{jj'})^{q_{jj'}} (q_{jj'}!)^{-1} \quad (3.63)$$

$$= \prod_{j=1}^J \Gamma(n_{j\cdot})^{-1} u_j^{-1} (1+u_j)^{-\alpha} \left(\frac{u_j}{1+u_j} \right)^{n_{j\cdot}+q_{j\cdot}} \quad (3.64)$$

$$\times \prod_{j'=1}^J \frac{\Gamma(\alpha\beta_{j'} + n_{jj'} + q_{jj'})}{\Gamma(\alpha\beta_{j'})} \phi_{jj'}^{n_{jj'}} (1 - \phi_{jj'})^{q_{jj'}} (q_{jj'}!)^{-1} \quad (3.65)$$

where we have used the fact that the β_j sum to 1. Therefore

$$p(\beta | z, u, Q, \alpha, \gamma, \theta) \propto \prod_{j=1}^J \beta_j^{\gamma_j-1} \prod_{j'=1}^J \frac{\Gamma(\alpha\beta_{j'} + n_{jj'} + q_{jj'})}{\Gamma(\alpha\beta_{j'})}. \quad (3.66)$$

Following (Teh et al., 2006), we can write the ratios of Gamma functions as polynomials in β_j , as

$$p(\beta | z, u, Q, \alpha, \gamma, \theta) \propto \prod_{j=1}^J \beta_j^{\gamma_j-1} \prod_{j'=1}^J \sum_{m_{jj'}=1}^{n_{jj'}} s(n_{jj'} + q_{jj'}, m_{jj'}) (\alpha\beta_{j'})^{m_{jj'}} \quad (3.67)$$

where $s(m, n)$ is an unsigned Stirling number of the first kind, which is used to represent the number of permutations of n elements such that there are m distinct cycles.

This admits an augmented data representation, where we introduce a random matrix $M = (m_{jj'})_{1 \leq j, j' \leq J}$, whose entries are conditionally independent given β , Q and z , with

$$p(m_{jj'} = m | \beta_{j'}, \alpha, n_{jj'}, q_{jj'}) = \frac{s(n_{jj'} + q_{jj'}, m) \alpha^m \beta_{j'}^m}{\sum_{m'=0}^{n_{jj'}+q_{jj'}} s(n_{jj'} + q_{jj'}, m') \alpha^{m'} \beta_{j'}^{m'}} \quad (3.68)$$

for integer m ranging between 0 and $n_{jj'} + q_{jj'}$. Note that $s(n, 0) = 0$ if $n > 0$, $s(0, 0) = 1$, $s(0, m) = 0$ if $m > 0$, and we have the recurrence relation $s(n+1, m) =$

$ns(n, m) + s(n, m - 1)$, and so we could compute each of these coefficients explicitly; however, it is typically simpler and more computationally efficient to sample from this distribution than it is to enumerate its probabilities

For each $m_{jj'}$ we simply draw $n_{jj'}$ assignments of customers to tables according to the Chinese Restaurant Process and set $m_{jj'}$ to be the number of distinct tables realized; that is, assign the first customer to a table, setting $m_{jj'}$ to 1, and then, after n customers are assigned, assign the $n + 1$ th customer to a new table with probability $\alpha\beta_{j'}/(n + \alpha\beta_{j'})$, in which case we increment $m_{jj'}$, and to an existing table with probability $n/(n + \alpha)$, in which case we do not increment $m_{jj'}$.

Then, we have joint distribution

$$p(\beta, M | z, u, Q, \alpha, \gamma, \ell) \propto \prod_{j=1}^J \beta_j^{\frac{\gamma}{J}-1} \prod_{j'=1}^J s(n_{jj'} + q_{jj'}, m_{jj'}) \alpha^{m_{jj'}} \beta_{j'}^{m_{jj'}} \quad (3.69)$$

which yields (3.67) when marginalized over M . Again discarding constants in β and regrouping yields

$$p(\beta | M, z, u, \theta, \alpha, \gamma) \propto \prod_{j'=1}^J \beta_{j'}^{\frac{\gamma}{J} + m_{\cdot, j'} - 1} \quad (3.70)$$

which is Dirichlet:

$$\beta | M, \gamma \sim \text{Dirichlet}\left(\frac{\gamma}{J} + m_{\cdot, 1}, \dots, \frac{\gamma}{J} + m_{\cdot, J}\right) \quad (3.71)$$

Sampling α and γ Assume that α and γ have Gamma priors, with

$$p(\alpha) = \frac{b_\alpha^{a_\alpha}}{\Gamma(a_\alpha)} \alpha^{a_\alpha-1} \exp(-b_\alpha \alpha) \quad (3.72)$$

$$p(\gamma) = \frac{b_\gamma^{a_\gamma}}{\Gamma(a_\gamma)} \gamma^{a_\gamma-1} \exp(-b_\gamma \gamma) \quad (3.73)$$

Having integrated out π , we have

$$p(\beta, z, u, Q, M \mid \alpha, \gamma, \theta) = \frac{\Gamma(\gamma)}{\Gamma(\frac{\gamma}{J})^J} \alpha^{m..} \prod_{j=1}^J \beta_j^{\frac{\gamma}{J} + m_{..j} - 1} \Gamma(n_{j.})^{-1} u_j^{-1} (1 + u_j)^{-\alpha} \left(\frac{u_j}{1 + u_j} \right)^{n_{j.} + q_{j.}} \quad (3.74)$$

$$\times \prod_{j'=1}^J s(n_{jj'} + q_{jj'}, m_{jj'}) \phi_{jj'}^{n_{jj'}} (1 - \phi_{jj'})^{q_{jj'}} (q_{jj'}!)^{-1} \quad (3.75)$$

We can also integrate out β , to yield

$$p(z, u, Q, M \mid \alpha, \gamma, \theta) = \alpha^{m..} e^{-\sum_{j''} \log(1 + u_{j''}) \alpha} \frac{\Gamma(\gamma)}{\Gamma(\gamma + m_{..})} \quad (3.76)$$

$$\times \prod_j \frac{\Gamma(\frac{\gamma}{J} + m_{..j})}{\Gamma(\frac{\gamma}{J}) \Gamma(n_{j.})} u_j^{-1} \left(\frac{u_j}{1 + u_j} \right)^{n_{j.} + q_{j.}} \quad (3.77)$$

$$\times \prod_{j'=1}^J s(n_{jj'} + q_{jj'}, m_{jj'}) \phi_{jj'}^{n_{jj'}} (1 - \phi_{jj'})^{q_{jj'}} (q_{jj'}!)^{-1} \quad (3.78)$$

demonstrating that α and γ are independent given θ and the augmented data, with

$$p(\alpha \mid z, u, Q, M, \theta) \propto \alpha^{a_\alpha + m_{..}} \exp(-(b_\alpha + \sum_j \log(1 + u_j)) \alpha) \quad (3.79)$$

and

$$p(\gamma \mid z, u, Q, M, \theta) \propto \gamma^{a_\gamma - 1} \exp(-b_\gamma \gamma) \frac{\Gamma(\gamma) \prod_{j=1}^J \Gamma(\frac{\gamma}{J} + m_{..j})}{\Gamma(\frac{\gamma}{J})^J \Gamma(\gamma + m_{..})} \quad (3.80)$$

So we see that

$$\alpha \mid z, u, Q, M, \theta \sim \text{Gamma}(a_\alpha + m_{..}, b_\alpha + \sum_j \log(1 + u_j)) \quad (3.81)$$

To sample γ , we introduce a new set of auxiliary variables, $r = (r_1, \dots, r_J)$ and t with the following distributions:

$$p(r_{j'} = r \mid m_{..j'}, \gamma) = \frac{\Gamma(\frac{\gamma}{J})}{\Gamma(\frac{\gamma}{J} + m_{..j'})} s(m_{..j'}, r) \left(\frac{\gamma}{J} \right)^r \quad r = 1, \dots, m_{..j'} \quad (3.82)$$

$$p(t \mid m_{..}, \gamma) = \frac{\Gamma(\gamma + m_{..})}{\Gamma(\gamma) \Gamma(m_{..})} t^{\gamma-1} (1-t)^{m_{..}-1} \quad t \in (0, 1) \quad (3.83)$$

so that

$$p(\gamma, r, t \mid M) \propto \gamma^{a_\gamma-1} \exp(-b_\gamma \gamma) t^{\gamma-1} (1-t)^{m_{..}-1} \prod_{j'=1}^J s(m_{.j'}, r_{j'}) \left(\frac{\gamma}{J}\right)^{r_{j'}} \quad (3.84)$$

and

$$p(\gamma \mid r, t) \propto \gamma^{a_\gamma+r_{..}-1} \exp(-(b_\gamma - \log(t))\gamma), \quad (3.85)$$

which is to say

$$\gamma \mid r, t, z, u, Q, M, \theta \sim \text{Gamma}(a_\gamma + r_{..}, b_\gamma - \log(t)) \quad (3.86)$$

Summary I have made the following additional assumptions about the generative model in this section:

$$\gamma \sim \text{Gamma}(a_\gamma, b_\gamma) \quad \alpha \sim \text{Gamma}(a_\alpha, b_\alpha) \quad (3.87)$$

The joint conditional over γ , α , β and π given z , u , Q , M , r , t and θ factors as

$$p(\gamma, \alpha, \beta, \pi \mid z, u, Q, r, t, \theta) = p(\gamma \mid r, t) p(\alpha \mid u, M) p(\beta \mid \gamma, M) p(\pi \mid \alpha, \beta, z, u, Q) \quad (3.88)$$

where

$$\gamma \mid r, t \sim \text{Gamma}(a_\gamma + r_{..}, b_\gamma - \log(t)) \quad (3.89)$$

$$\alpha \mid u, M \sim \text{Gamma}(a_\alpha + m_{..}, b_\alpha + \sum_j \log(1 + u_j)) \quad (3.90)$$

$$\beta \mid \gamma, M \sim \text{Dirichlet}\left(\frac{\gamma}{J} + m_{.1}, \dots, \frac{\gamma}{J} + m_{.J}\right) \quad (3.91)$$

$$\pi_{jj'} \mid \alpha, \beta_{j'}, z, u, Q \stackrel{\text{ind}}{\sim} \text{Gamma}(\alpha \beta_{j'} + n_{jj'} + q_{jj'}, 1 + u_j) \quad (3.92)$$

3.4.2. Sampling z and the auxiliary variables

The hidden state sequence, z , is sampled jointly with the auxiliary variables, which consist of u , M , Q , r and t . The joint conditional distribution of these variables is

defined directly by the generative model:

$$p(z, u, Q, M, r, t \mid \pi, \beta, \alpha, \gamma, \ell) = p(z \mid \pi, \theta) p(u \mid z, \pi, \ell) p(Q \mid u, \pi, \ell) p(M \mid z, Q, \alpha, \beta) \quad (3.93)$$

$$\times p(r \mid \gamma, M) p(t \mid \gamma, M) \quad (3.94)$$

Since we are representing the transition matrix explicitly, we can sample the entire sequence z at once with the forward-backward algorithm, as in an ordinary HMM (or, if we are employing the HSMM variant described in Sec. 3.3.1, then we can use the modified message passing scheme for HSMMs described by Johnson and Willsky (2013)). Having done this, we can sample u , Q , M , r and t from their forward distributions. To summarize, we have

$$u_j \mid z, \pi, \ell \stackrel{ind}{\sim} \text{Gamma}(n_{j\cdot}, \sum_{j'} \pi_{jj'} \phi_{jj'}) \quad (3.95)$$

$$q_{jj'} \mid u_j, \pi_{jj'}, \phi_{jj'} \stackrel{ind}{\sim} \text{Poisson}(u_j(1 - \phi_{jj'})\pi_{jj'}) \quad (3.96)$$

$$m_{jj'} \mid n_{jj'}, q_{jj'}, \beta_{j'}, \alpha \stackrel{ind}{\sim} \frac{\Gamma(\alpha\beta_j)}{\Gamma(\alpha\beta_j + n_{jj'} + q_{jj'})} \sum_{m=1}^{n_{jj'} + q_{jj'}} s(n_{jj'} + q_{jj'}, m) \alpha^m \beta_{j'}^m \delta_m \quad (3.97)$$

$$r_j \mid m_{\cdot j}, \gamma \stackrel{ind}{\sim} \frac{\Gamma(\frac{\gamma}{J})}{\Gamma(\frac{\gamma}{J} + m_{\cdot j})} \sum_{r=1}^{m_{\cdot j}} s(m_{\cdot j}, r) \left(\frac{\gamma}{J}\right)^r \delta_r \quad (3.98)$$

$$t \mid \gamma, M \sim \text{Beta}(\gamma, m_{\cdot\cdot}) \quad (3.99)$$

3.4.3. Sampling state and emission parameters

The state parameters, θ , influence the transition matrix, π and the auxiliary vector q through the similarity matrix matrix ϕ , and also control the emission distributions. We have likelihood factors

$$p(z, Q \mid \theta) \propto \prod_j \prod_{j'} \phi_{jj'}^{n_{jj'}} (1 - \phi_{jj'})^{q_{jj'}} \quad (3.100)$$

$$p(y \mid z, \theta) = \prod_{t=1}^T f(y_t; \theta_{z_t}) \quad (3.101)$$

where, recall, $n_{jj'}$ counts the number of times that the state sequence transitions from state j to state j' , θ_j is the part of ℓ_j that governs the emission distribution for state j , and proportionality is with respect to variation in θ .

The parameter space for the hidden states, the associated prior H on θ , and the similarity function ϕ , is application-specific, thus I now turn to three individual applications with qualitatively different state spaces and emission distributions, derive inference methods for each, and present experiments on synthetic and real data.

3.5. Use Cases

In Chapter 4, I describe an application of the HaMMLeT model to a synthetic dataset designed to mimic a speaker diarization or blind source separation task. Here, each latent state corresponds to a description of which of several sound sources is active at a given time step, and where the observed data is a set of signals picked up from several microphones distributed around the room, each of which picks up all of the sound sources with varying sensitivities, and the goal is to determine who is speaking when. In this application the latent states are represented as binary vectors, similarity is based on Hamming distance between binary vectors, and the emission model is a multivariate linear-Gaussian model.

In Chapter 5 I describe a power disaggregation application, in which the observation sequence consists the amount of power used by a house at various times throughout a day, and the goal is to separate the aggregated power signal into several signals corresponding to a set of appliances (oven, air conditioning, lighting, etc.). Unlike the cocktail party setting, each latent signal may have more than two levels, and so in addition to inferring what appliances are on when, the model needs to discover how much power each appliance uses in each of its latent states. Here, latent states are vectors where each entry is a categorical label, the number of categories per channel is unknown, the similarity model is again based on Hamming distance

between binary vectors, and the emission model is linear-Gaussian.

Finally, in Chapter 6 I describe an application to unsupervised learning of musical structure. The data is a sequence of chords in a musical composition, represented as a single symbol, and the goal is to infer a set of chord equivalence classes and a transition model between equivalence classes. Unlike in the previous applications, similarity between states is modeled separately from the emission distributions, so that $\ell_j = (\eta_j, \theta_j)$, where ϕ is based on the η_j part of the state parameters, and θ_j is a disjoint set of parameters governing the emission distribution.

4. BINARY VECTOR STATES: SPEAKER DIARIZATION

4.1. Binary State Vectors

I consider here the case where a state consists of a finite D -dimensional binary vector, ℓ_j , the emission parameter vector θ_j is all of ℓ_j , and the emission distribution is based on a noisy linear transformation of the state vector according to a $D \times K$ weight matrix W , so that each K -dimensional observation is $\mathcal{N}(W^\top \theta_j, \Sigma)$. For the experiments discussed in this chapter, I will assume that Σ does not depend on j , but this assumption is easily relaxed if appropriate. For finite-length binary vector states, the set of possible states is finite, and so on its face it may seem that a nonparametric model is unnecessary. However, if D is reasonably large, it is likely that most of the 2^D possible states are vanishingly unlikely (and, in fact, the number of observations may well be less than 2^D), and so we would like a model that encourages the selection of a sparse set of states. Moreover, there could be more than one state with the same emission parameters, but with different transition dynamics.

We require a similarity function, $\phi(\theta_j, \theta_{j'})$, which varies between 0 to 1, and is equal to 1 if and only if $\theta_j = \theta_{j'}$. A natural choice in this setting is the Laplacian kernel:

$$\phi_{jj'} = \phi(\theta_j, \theta_{j'}) = \exp(-\lambda \Delta_{jj'}) \quad (4.1)$$

where $\Delta_{jj'd} = |\theta_{jd} - \theta_{j'd}|$, $\Delta_{jj'} = \sum_{d=1}^D \Delta_{jj'd}$ is the Hamming distance between θ_j and $\theta_{j'}$, and $\lambda \geq 0$ (if $\lambda = 0$, the $\phi_{jj'}$ are all 1, and so do not have any influence, reducing the model to an ordinary HDP-HMM).

Before describing individual experiments, I describe the additional inference steps needed for these variables.

4.1.1. Additional Inference Steps

Sampling θ We put independent Beta-Bernoulli priors on each coordinate of θ . Each coordinate θ_{jd} can be sampled conditioned on all the others and the coordinate-wise prior means, $\{\mu_d\}$, which we sample in turn conditioned on the θ s. Since individual coordinates are binary, each one has a Bernoulli posterior, whose parameter we now derive by computing the posterior log-odds, that is we will derive

$$\log \frac{p(\theta_{jd} = 1 \mid \theta \setminus \theta_{jd}, z, Q, Y)}{p(\theta_{jd} = 0 \mid \theta \setminus \theta_{jd}, z, Q, Y)} \quad (4.2)$$

Since θ_{jd} affects both the similarity matrix and the emission distribution, the posterior log-odds has three components: the prior log odds, the likelihood due to the successful and failed transitions (specifically the outgoing transition counts from state j , $n_{jj'}$ and $q_{jj'}$, as well as the incoming transition counts to state j , $n_{j'j}$ and $q_{j'j}$, across all j'), and the likelihood due to the data, y , for those t such that $z_t = j$.

The prior log odds is simply $\log(\mu_d/(1 - \mu_d))$. Next, we compute the likelihood component due to the transition attempts, that is

$$\log \frac{p(z, Q \mid \theta_{jd} = 1, \theta \setminus \theta_{jd}, \lambda)}{p(z, Q \mid \theta_{jd} = 0, \theta \setminus \theta_{jd}, \lambda)}$$

Let

$$\phi_{jj'-d} = \exp(-\lambda(\Delta_{jj'} - \Delta_{jj'd})) \quad (4.3)$$

so that $\phi_{jj'} = \phi_{jj'-d} e^{-\lambda \Delta_{jj'd}}$.

Since the matrix ϕ is assumed to be symmetric, we have

$$\frac{p(z, Q \mid \theta_{jd} = 1, \theta \setminus \theta_{jd})}{p(z, Q \mid \theta_{jd} = 0, \theta \setminus \theta_{jd})} \propto \prod_{j' \neq j} \frac{e^{-\lambda(n_{jj'} + n_{j'j})|1 - \theta_{j'd}|} (1 - \phi_{jj'-d} e^{-\lambda|1 - \theta_{j'd}|})^{q_{jj'} + q_{j'j}}}{e^{-\lambda(n_{jj'} + n_{j'j})|\theta_{j'd}|} (1 - \phi_{jj'-d} e^{-\lambda|\theta_{j'd}|})^{q_{jj'} + q_{j'j}}} \quad (4.4)$$

$$= e^{-\lambda(c_{jd0} - c_{jd1})} \prod_{j' \neq j} \left(\frac{1 - \phi_{jj'-d} e^{-\lambda}}{1 - \phi_{jj'-d}} \right)^{(-1)^{\theta_{j'd}}(q_{jj'} + q_{j'j})} \quad (4.5)$$

where c_{jd0} and c_{jd1} are the number of successful jumps to or from state j , to or from states with a 0 or 1, respectively, in position d . That is,

$$c_{jd0} = \sum_{\{j' \mid \theta_{j'd}=0\}} n_{jj'} + n_{j'j} \quad c_{jd1} = \sum_{\{j' \mid \theta_{j'd}=1\}} n_{jj'} + n_{j'j} \quad (4.6)$$

We assume the observed data Y consists of a $T \times K$ matrix, where the t th row $y_t = (y_{t1}, \dots, y_{tK})$ is a K -dimensional feature vector associated with time t , and let W be a $D \times K$ weight matrix with k th column w_k that maps D dimensional binary state vectors to means in the K -dimensional observation space. Then the part of the likelihood function from time t is

$$p(y_t \mid z_t, \theta_{z_t}) = \mathcal{N}(y_t \mid W^\top \theta_{z_t}, \Sigma) \quad (4.7)$$

In the experiments discussed here, Σ is assumed diagonal with entries σ_k^2 , $k = 1, \dots, K$, and does not depend on j . Under this assumption, we have

$$p(y_t \mid W^\top \theta_{z_t}) = \prod_{k=1}^K \mathcal{N}(y_t \mid w_k^\top \theta_{z_t}, \sigma_k^2). \quad (4.8)$$

The coordinate θ_{jd} affects the likelihood only through the mean, $W^\top \theta_j$. Expanding this mean to isolate the effect of θ_{jd} , we have

$$w_k^\top \theta_j = \sum_d w_{dk} \theta_{jd}$$

and so we see that setting $\theta_{jd} = 1$ vs 0 shifts the k th coordinate of the mean, $W^\top \theta_j$, by w_{dk} . Denote by θ^* the matrix whose t th row consists of the binary state vector θ_{z_t} , and let $X^* = W^\top \theta^*$. As shorthand, write $x_{tk}^{(-d)}$ to be the (t, k) entry in X^* assuming that $\theta_{jd} = 0$. Then the corresponding value if $\theta_{jd} = 1$ is $x_{tk}^{(-d)} + w_{dk}$. Since the y_t are conditionally independent given the state sequence, z , we can write the log likelihood ratio for $\theta_{jd} = 1$ versus $\theta_{jd} = 0$ as

$$\sum_{t: z_t=j} \sum_{k=1}^K \log \left(\frac{\mathcal{N}(y_{tk} \mid x_{tk}^{(-d)} + w_{dk}, \sigma_k^2)}{\mathcal{N}(y_{tk} \mid x_{tk}^{(-d)}, \sigma_k^2)} \right) = \sum_{t: z_t=j} \sum_{k=1}^K -\frac{w_{dk}}{\sigma_k^2} (y_{tk} - x_{tk}^{(-d)} + \frac{1}{2} w_{dk}) \quad (4.9)$$

All together, we have

$$\log \left(\frac{p(\theta_{jd} = 1 | Y, z, Q, \theta \setminus \theta_{jd})}{p(\theta_{jd} = 0 | Y, z, Q, \theta \setminus \theta_{jd})} \right) \quad (4.10)$$

$$= \log \left(\frac{p(\theta_{jd} = 1)p(z, Q | \theta_{jd} = 1, \theta \setminus \theta_{jd})p(Y | z, \theta_{jd} = 1, \theta \setminus \theta_{jd})}{p(\theta_{jd} = 0)p(z, Q | \theta_{jd} = 0, \theta \setminus \theta_{jd})p(Y | z, \theta_{jd} = 0, \theta \setminus \theta_{jd})} \right) \quad (4.11)$$

$$= \log \left(\frac{\mu_d}{1 - \mu_d} \right) \quad (4.12)$$

$$+ (c_{jd1} - c_{jd0})\lambda + \sum_{j' \neq j} (-1)^{\theta_{j'd}} (q_{jj'} + q_{j'j}) \log \left(\frac{1 - \phi_{jj'}^{(-d)} e^{-\lambda}}{1 - \phi_{jj'}^{(-d)}} \right) \quad (4.13)$$

$$+ \sum_{t: z_t = j} \sum_{k=1}^K -\frac{w_{dk}}{\sigma_k^2} (y_{tk} - x_{tk}^{(-d)} + \frac{1}{2} w_{dk}) \quad (4.14)$$

Sampling μ Sampling the μ_d is straightforward with a Beta prior. Suppose

$$\mu_d \stackrel{ind}{\sim} \text{Beta}(a_\mu, b_\mu) \quad (4.15)$$

Then, conditioned on θ the μ_d are independent with

$$\mu_d | \theta \sim \text{Beta}(a_\mu + \sum_j \theta_{jd}, b_\mu + \sum_j (1 - \theta_{jd})) \quad (4.16)$$

Sampling W and Σ Conditioned on the state matrix θ and the data matrix Y , the weight matrix W can be sampled as well using standard methods for Bayesian linear regression. We place a zero mean Normal prior on each element of W (including a row of intercept terms), resulting in a multivariate Normal posterior for each column. For the experiments reported below, we constrain Σ to be a diagonal matrix, and place an Inverse Gamma prior on the variances, resulting in conjugate updates.

Define the prior:

$$p(W) = \prod_{k=1}^K \prod_{d=1}^D \mathcal{N}(w_{dk} | 0, \sigma_0^2). \quad (4.17)$$

The likelihood is

$$p(Y | W, \theta) = \prod_{t=1}^T \prod_{k=1}^K p(y_{tk}; x_{tk}) = \prod_{t=1}^T \prod_{k=1}^K \mathcal{N}(y_{tk} | x_{tk}, \sigma_k^2) \quad (4.18)$$

Then it is a standard result from Bayesian linear modeling that

$$p(W | \theta, Y) = \prod_{k=1}^K \mathcal{N}((\sigma_0^{-2} \mathbf{I} + \Sigma^{-1} \theta^{*\top} \theta^*)^{-1} \Sigma^{-1} \theta^{*\top} Y, (\sigma_0^{-2} \mathbf{I} + \Sigma^{-1} \theta^{*\top} \theta^*)^{-1}) \quad (4.19)$$

Sampling λ The Laplacian kernel ϕ is defined as $\phi(\theta_j, \theta_{j'}) = e^{-\lambda d(\theta_j, \theta_{j'})}$, where in our case d is Hamming distance. The parameter λ governs the connection between θ and ϕ . Writing (3.100) in terms of λ and the distance matrix Δ gives the likelihood

$$p(z, Q | \lambda, \theta) \propto \prod_j \prod_{j'} e^{-\lambda \Delta_{jj'} n_{jj'}} (1 - e^{-\lambda \Delta_{jj'}})^{q_{jj'}} \quad (4.20)$$

We put an **Exponential**(b_λ) prior on λ , which yields a posterior density

$$p(\lambda | z, Q, \theta) \propto e^{-(b_\lambda + \sum_j \sum_{j'} \Delta_{jj'} n_{jj'}) \lambda} \prod_j \prod_{j'} (1 - e^{-\lambda \Delta_{jj'}})^{q_{jj'}} \quad (4.21)$$

This density is log-concave, with

$$-\frac{d^2 \log(p(\lambda | z, Q, \theta))}{d\lambda^2} = \sum_{\{(j,j') | \Delta_{jj'} > 0\}} \frac{\Delta_{jj'}^2 q_{jj'} e^{\lambda \Delta_{jj'}}}{(e^{\lambda \Delta_{jj'}} - 1)^2} > 0 \quad (4.22)$$

and so we can use Adaptive Rejection Sampling (Gilks and Wild, 1992) to sample from it. To do this we need to compute the log density and its first derivative, which are given by

$$h(\lambda) = -(b_\lambda + \sum_{\{(j,j') | \Delta_{jj'} > 0\}} \Delta_{jj'} n_{jj'}) \lambda + \sum_{\{(j,j') | \Delta_{jj'} > 0\}} q_{jj'} \log(1 - e^{-\lambda \Delta_{jj'}}) \quad (4.23)$$

$$h'(\lambda) = -(b_\lambda + \sum_{\{(j,j') | \Delta_{jj'} > 0\}} \Delta_{jj'} n_{jj'}) + \sum_{\{(j,j') | \Delta_{jj'} > 0\}} \frac{q_{jj'} \Delta_{jj'}}{e^{\lambda \Delta_{jj'}} - 1} \quad (4.24)$$

4.1.2. Summary

I have made the following assumptions about the representation of the hidden states and observed data in this subsection: (1) States are distinguished via a D -dimensional binary parameter vector, θ_j (2) the similarity ϕ between states j and j' is the Laplacian

kernel with respect to Hamming distance between the respective θ vectors, and has a decay parameter λ , and (3) Y consists of K real-valued features associated with each time step t . In addition, the following priors are applied:

$$\mu_d \stackrel{i.i.d.}{\sim} \text{Beta}(a_\mu, b_\mu) \quad (4.25)$$

$$\lambda \sim \text{Exponential}(b_\lambda) \quad (4.26)$$

$$\theta_{jd} | \mu \stackrel{ind}{\sim} \text{Bernoulli}(\mu_d) \quad (4.27)$$

$$w_{dk} \stackrel{i.i.d.}{\sim} \mathcal{N}(0, \sigma_0^2) \quad (4.28)$$

$$y_{tk} | W, z, \theta \stackrel{ind}{\sim} \mathcal{N}(w_k^T \theta_{z_t}, \Sigma) \quad (4.29)$$

I introduce Gibbs blocks corresponding to (1) each θ_{jd} individually, (2) the vector μ , (3) the decay parameter λ , and (4) the weight matrix W . We have conditional posterior distributions

$$\theta_{jd} | \theta \setminus \theta_{jd}, z, Q, \mu, \lambda, W, Y^* \sim \text{Bernoulli} \left(\frac{e^{\zeta_{jd}}}{1 + e^{\zeta_{jd}}} \right) \quad (4.30)$$

$$\mu_d | \theta, \dots \stackrel{ind}{\sim} \text{Beta}(a_\mu + \sum_j \theta_{jd}, b_\mu + \sum_j (1 - \theta_{jd})) \quad (4.31)$$

$$p(\lambda | z, Q, \theta, \dots) \propto e^{-(b_\lambda + \sum_j \sum_{j'} \Delta_{jj'} n_{jj'}) \lambda} \prod_j \prod_{j'} (1 - e^{-\lambda \Delta_{jj'}})^{q_{jj'}} \quad (4.32)$$

$$W | \theta, Y, z, \dots \stackrel{ind}{\sim} \mathcal{N}((\sigma_0^{-2} \mathbf{I} + \Sigma^{-1} \theta^{*T} \theta^*)^{-1} \Sigma^{-1} \theta^{*T} Y, (\sigma_0^{-2} \mathbf{I} + \Sigma^{-1} \theta^{*T} \theta^*)^{-1}) \quad (4.33)$$

where $\Delta_{jj'} = \|\theta_j - \theta_{j'}\|_{L_1}$ and

$$\begin{aligned} \zeta_{jd} = & \log \left(\frac{\mu_d}{1 - \mu_d} \right) + (c_{jd1} - c_{jd0}) \lambda + \sum_{j' \neq j} (-1)^{\theta_{j'd}} (q_{jj'} + q_{j'j}) \log \left(\frac{1 - \phi_{jj'}^{(-d)} e^{-\lambda}}{1 - \phi_{jj'}^{(-d)}} \right) \\ & - \sum_{\{t | z_t = j\}} \sum_{k=1}^K \frac{w_{dk}}{\sigma_k^2} (y_{tk}^* - x_{tk}^{(-d)} + \frac{1}{2} w_{dk}) \end{aligned} \quad (4.34)$$

All distributions can be sampled from directly except for λ , which requires Adaptive

Rejection Sampling, with the equations

$$h(\lambda) = -(b_\lambda + \sum_{\{(j,j') \mid \Delta_{jj'} > 0\}} \Delta_{jj'} n_{jj'}) \lambda + \sum_{\{(j,j') \mid \Delta_{jj'} > 0\}} q_{jj'} \log(1 - e^{-\lambda \Delta_{jj'}}) \quad (4.35)$$

$$h'(\lambda) = -(b_\lambda + \sum_{\{(j,j') \mid \Delta_{jj'} > 0\}} \Delta_{jj'} n_{jj'}) + \sum_{\{(j,j') \mid \Delta_{jj'} > 0\}} \frac{q_{jj'} \Delta_{jj'}}{e^{\lambda \Delta_{jj'}} - 1} \quad (4.36)$$

4.2. “Cocktail Party” Data

To evaluate the model, we created synthetic data based on a speaker diarization/blind source separation task known as the “cocktail party” problem. In such a task, the observable data consists of one or more audio signals, discretized into T segments, where the signal at time t is assumed to arise from multiple speakers, some subset of whom are speaking. The inference problem is to recover which speakers are speaking when. We further assume that speakers are grouped into conversations, and take turns speaking within conversation. In such a task, there are naively 2^S possible states, where each state indicates which of S speakers is speaking at a particular time step. However, due to the conversational grouping, if at most one speaker in a conversation can be speaking at any given time (a natural assumption of within-group turn-taking), the state space is constrained, with only $\prod_c (s_c + 1)$ states possible (with perhaps a small probability of other combinations occurring due to occasional overlap), where s_c is the number of speakers in conversation c .

For the synthetic cocktail party data, the turn sequence within each conversation is generated using a Hidden Semi-Markov Model (HSMM) with Poisson distributed state-durations, in which the state sequence had s_c states, with pauses with shorter Poisson duration inserted between each “sentence”. The states within conversations are then mapped to a binary vector with s_c entries, where silence is represented by all zeroes, and speaker s speaking corresponds to a 1 in position s . The binary vectors are concatenated across conversations to yield latent states consisting of length S binary vectors. To simulate speakers being recorded by K microphones, weights mapping

speakers to microphones were generated independently from a $U(0, 1)$ distribution, resulting in a $D \times K$ weight matrix, W . A constant “background sound level” parameter for each microphone was added as well, also $U(0, 1)$, and independent $\mathcal{N}(0, \sigma_k^2)$ noise was added to each time step at microphone k .

Transition and emission parameters were generated from conjugate priors to the Poisson HSMM. The data set consisted of four conversations of four speakers each, and 12 microphones, so that $D = 16$, and $K = 12$. There are therefore $2^{16} = 65536$ possible binary vector-valued states, but only $(4+1)^4 = 625$, less than 1%, can actually occur. The noise variance σ_k^2 was set to a constant of 1/10 for all $k = 1, \dots, K$.

We attempted to infer the states from the data using three models: (1) a binary-state Factorial HMM, in which the individual binary speaker sequences are modeled as independent a priori, (2) an ordinary HDP-HMM without local transitions, where the latent states are binary vectors, and (3) our HDP-HMM-LT model. The same linear-Gaussian emission model was employed for all three models; the only difference was the prior on the sequence of binary state vectors.

To simplify interpretation of the results, the weight matrix was fixed to the true value (this makes the latent dimensions identifiable and makes comparisons between inferred and ground truth state matrices meaningful). Each model was evaluated at each Gibbs iteration using the Hamming distance between inferred and ground truth state matrices, as well as the F1 score, which is the harmonic mean between the precision (the proportion of the 1s in the inferred state matrix that were actually 1 in the ground truth) and recall (the proportion of the 1s in the ground truth state matrix that were correctly classified as 1 by the model).

The results for the three models are in Figure 4.1. The LT model outperforms the other two on all measures on all datasets.

We also plot the inferred decay rate λ for the HDP-HMM-LT model. The LT model settles on a non-negligible λ value for this data, suggesting that the local transition structure explains the data well. It also uses more components than the

non-LT model, perhaps owing to the fact that the weaker transition prior of the non-LT model is more likely to explain nearby similar observations as a single persisting state, whereas the LT model places a higher probability on transitioning to a new state with a similar latent vector.

4.3. Synthetic Data Without Local Transitions

As a “sanity check” on the model, data was also generated directly from an ordinary HDP-HMM, with no local transition property, in order to investigate the performance of our model in a case where the data did not have the key property that its prior equipped it to discover. The results are in Fig. 4.2. During iterations when the λ parameter is large, the LT model has worse performance than the non-LT model on this data, as its bias toward local transitions is not helpful; however, the λ parameter settles near zero as sampling continues, and the model learns that a preference for local transitions does not help to explain the data. Note that when $\lambda = 0$, the HDP-HMM-LT reduces to an ordinary HDP-HMM. Unlike on the cocktail party data, the LT model does not use more states when the data does not have the LT property.




Figure 4.1: (a-b) Accuracy and F1 scores for the HDP-HMM-LT, standard HDP-HMM, and Binary Factorial HMM on the Cocktail Party Data. Metrics are averaged over 10 Gibbs runs on each model, with error bars representing a 99% confidence interval for the mean per iteration. The first 100 iterations are not shown. (c) Learned similarity parameter, λ , for the LT model, (d) Number of distinct states used by HDP-HMM and HDP-HMM-LT. The first 100 iterations are excluded.

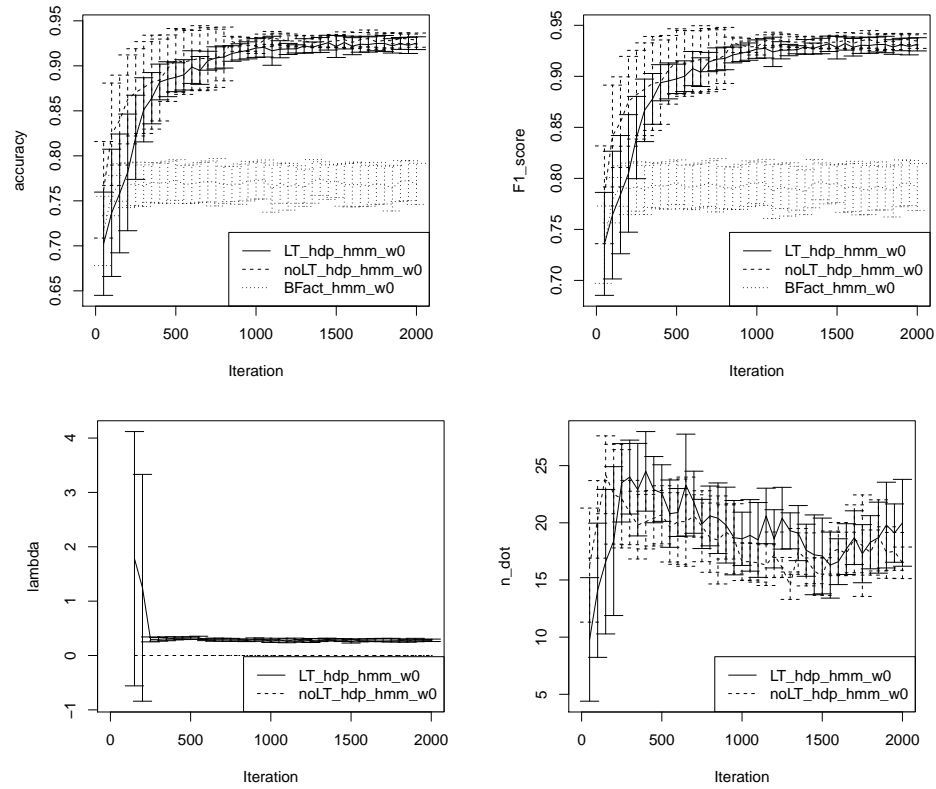


Figure 4.2: (a-b) Accuracy and F1 for the three models on data generated from an HDP-HMM without local transitions, (c) Learned similarity parameter, λ , for the LT model, (d) Number of states used by the HDP-HMM and HDP-HMM-LT. The first 100 iterations are omitted.

5. CATEGORICAL VECTOR STATES: POWER DISAGGREGATION

5.1. Generalizing to Categorical-Valued θ

For the experiment described in this chapter, we relax the assumption made in Chapter 4 that the θ_j are binary vectors and instead allow each θ_{jd} to take on one in an arbitrary set of discrete values — i.e., to be categorically distributed rather than Bernoulli distributed; however the similarity between two states will again be defined in terms of distance between θ_j vectors that also inform the emission distribution via a linear-Gaussian model. As a result, much of the additional inference steps described for the cocktail party experiment carry over here. The exception of course is inferences about θ .

In this chapter, I first define representations and priors for θ and the weight matrix, W for the general case of categorical-vector-valued θ . Then I derive the necessary Gibbs steps. Finally, I describe an application to power disaggregation, where the observations consists of an aggregated energy-use time series from several houses, the latent state vector describes the energy being used by each of several channels (e.g., lighting, refrigerator, washer/dryer), such that individual channels are in one of several discrete states at a given time. For example, the refrigerator channel might cycle between an “off” state, a low-power standby state, and a high-power “active cooling” state. Inference consists of discovering the set of discrete states for each channel, attributing a typical level of energy consumption for each state, and inferring what state each channel is in during each discretized time window in the data.

5.2. Priors and Representations in the Categorical State Variant

In place of Beta priors on each θ_{jd} , we use Chinese Restaurant Process priors, with concentration parameter $\alpha_d^{(\theta)}$. In the case that the number of categorical values is known in advance, this can be replaced by a Dirichlet prior, but I present the more flexible case here.

The weight matrix W must be expanded to allow for distinct weights associated with each possible value of the θ_{jd} . Rather than the matrix given by $(w_{dk})_{d=1,\dots,D,k=1,\dots,K}$, we now need a set of weights, (w_{sdk}) , where s indexes the categorical values that θ_{jd} can take. With a CRP prior, there is an unbounded number of s , though during inference, as in the direct assignment sampler for the HDP-HMM, we need only consider those values to which some state is currently assigned, plus one additional value representing a “new” state.

I will continue to assume that the ϕ function decays exponentially as a function of the number of component-wise differences (that is, Hamming distance) between a pair of states. Specifically, $\Delta_{jj'd}$ is zero if and only if $\theta_{jd} = \theta_{j'd}$ and is 1 otherwise, $\Delta_{jj'} = \sum_d \Delta_{jj'd}$, and $\phi(\theta_j \theta_{jj'}) = e^{-\lambda \Delta_{jj'}}$.

We can use a “dummy variable” representation of θ_j . Define S_d to be the number of realized states for dimension d and a 1 in position $\sum_{d' < d} S_{d'} + s$ indicates that $\theta_{jd} = s$. There will thus be D entries equal to 1, with the remaining entries equal to zero. We can then represent the weight matrix \mathbf{W} as stacked block matrix, where each block is $S_d \times K$, and there is one block for each d . In practice we only need to instantiate a new block when some θ_{jd} is assigned to a “new table” in the CRP metaphor, so that the dimension of \mathbf{W} is $\sum_d S_d \times K$. Then we have

$$y_t \sim \mathcal{N}(w^{(b)} + W^\top \theta_{z_t}, \Sigma) \quad (5.1)$$

where $w^{(b)}$ is a K -dimensional bias vector with a separate Normal prior, W is the weight matrix as defined just above, z_t is the state indicator for time t , and Σ is a $K \times K$ noise covariance matrix.

5.3. Adapting Posterior Inference for Categorical State Vectors

Sampling θ As before, the conditional posterior for θ_{jd} is proportional to the product of three terms: the prior (now a CRP), the likelihood of all successful and failed transitions to and from state j , and the likelihood of the observation sequence.

Under the CRP, the probability that $\theta_{jd} = s$ conditioned on the rest of θ (but not the data) is proportional to the number of other $j' \neq j$ such that $\theta_{j'd} = s$ where this count is positive; and proportional to $\alpha_j^{(\theta)}$ otherwise. Let

$$\tilde{n}_{ds}^{-j} = \sum_{j' \neq j} I(\theta_{j'd} = s), \quad s = 1, \dots, S_d \quad (5.2)$$

be these counts, where we assume that there are S_d distinct values taken by the θ_{jd} for a particular d . Then

$$p(\theta_{jd} = s \mid \theta_d^{-j}) \propto \begin{cases} \tilde{n}_{ds}^{-j} & s = 1, \dots, S_d \\ \alpha_d^{(\theta)} & s = S_d + 1 \end{cases} \quad (5.3)$$

The transition component of the likelihood is as in the binary case:

$$p(z, Q \mid \theta_{jd}, \theta_d^{-j}) \propto e^{-\lambda \sum_{j'} \Delta_{jj'}(n_{jj'} + n_{j'j})} \prod_{j' \neq j} (1 - e^{-\lambda \Delta_{jj'}(q_{jj'} + q_{j'j})}) \quad (5.4)$$

$$\propto e^{-\lambda \sum_{j'} I(\theta_{jd} \neq \theta_{j'd})(n_{jj'} + n_{j'j})} \prod_{j' \neq j} (1 - a \cdot e^{-\lambda I(\theta_{jd} \neq \theta_{j'd})(q_{jj'} + q_{j'j})}) \quad (5.5)$$

where a is a constant in θ_{jd} , defined as $e^{-\lambda \Delta_{jj'-d}(q_{jj'd} + q_{j'dj})}$. Taking a log yields

$$\log p(z, Q \mid \theta_{jd}, \theta_d^{-j}) = -\lambda \sum_{\{j': \theta_{jd} \neq \theta_{j'd}\}} (n_{jj'} + n_{j'j}) + \sum_{j' \neq j} \log(1 - a \cdot e^{-\lambda I(\theta_{jd} \neq \theta_{j'd})(q_{jj'} + q_{j'j})}) \quad (5.6)$$

The emission component of the likelihood is given for each t by

$$p(y_t \mid \theta_{zt}, W, \Sigma) \propto |\Sigma|^{-1/2} \exp\left(-\frac{1}{2}(y_t - w^{(b)} - W^\top \theta_{zt})^\top \Sigma^{-1}(y_t - w^{(b)} - W^\top \theta_{zt})\right) \quad (5.7)$$

Assuming a diagonal covariance matrix, isolating θ_{jd} , and taking a log, this be-

comes, for each k and t ,

$$\log p(y_{tk} | \theta_{z_t,d}, \theta_{z_t}^{-d}, w_k^{(b)}, \sigma_k^2) = -\frac{1}{2\sigma_k^2} \left(y_{tk} - w_k^{(b)} - \theta_{z_t} \mathbf{w}_k \right)^2 \quad (5.8)$$

$$\propto -\frac{1}{2\sigma_k^2} \left(\sum_{d'} w_{\theta_{z_t d'}, d', k} - (y_{tk} - w_k^{(b)}) \right)^2 \quad (5.9)$$

$$\propto -\frac{1}{2\sigma_k^2} \left(w_{\theta_{z_t d}, d, k} - (y_{tk} - w_k^{(b)} - \sum_{d' \neq d} w_{\theta_{z_t d'}, d', k}) \right)^2 \quad (5.10)$$

For a particular j and d , the full emission log likelihood is a sum of terms like the above, over all t such that $z_t = j$.

Sampling W Having expanded θ to a dummy variable representation, and having constructed a stacked block form of W , we can sample each column of W from a conditional posterior multivariate Normal just as in the binary case.

Sampling λ Since λ controls only the relationship between the distance matrix, Δ , and the similarity matrix, ϕ , upon conditioning on Δ and the jump attempts, λ is not sensitive to the way that Δ was computed; hence it can be sampled using Adaptive Rejection Sampling from exactly the same conditional distribution as before.

5.4. Power Disaggregation

The categorical model was tested using a subset of the Reference Energy Disaggregation Data set (REDD; Kolter and Johnson (2011)). The data consists of power consumption measurements from several time intervals from several electrical channels in several different homes, as well as an aggregated power consumption time series for each interval per home. An example interval is shown in Fig. 5.1.

Two key properties of this data are apparent in the figure. First, up to some small-scale variability, each channel visits a relatively small number of amplitudes, with some (such as the oven) alternating between an “on” state and an “off” state, and

others (such as the outlets and lighting channels) exhibiting more complex dynamics, presumably corresponding to various appliances or light fixtures turning on and off in different combinations. This property motivates the choice of a categorical state model. Second, there are clear correlations among the various channels: in this example, the two oven channels always go on and off together, albeit exhibiting different amplitudes when on, while two of the three lighting channels are highly but not perfectly correlated, while the third exhibits opposing behavior to the other two (perhaps arising from behavior).

Collectively, the dynamics exhibited here are similar to the cocktail party data described in Chapter 4. For one, since transitions between combinatorial states tend to be “local”, in that it is rare for more than one channel to change state at one time, we would expect the HaMMLeT model’s bias toward local state transitions to be beneficial, compared to the “vanilla” HDP-HMM which has no such locality bias. On the other hand, not all combinations of states occur in the data, and moreover, some state combinations are more or likely than the product of the individual channel probabilities would suggest, suggesting that the flexibility of a model such as HaMMLeT, in which there is a single latent state space over vector-valued states, might be better able to capture the correct transition probabilities than a model such as the Factorial HMM, in which the component chains evolve *a priori* independently.




Figure 5.1: A sample data interval from the REDD dataset Kolter and Johnson (2011). The top channel contains the total measured power in watts consumed by a home during a period of approximately 24 hours. Each timestep represents a 20 second intervals, during which the amplitude recorded is a median of the amplitudes in the original higher resolution data.

6. SEPARATE SIMILARITIES AND EMISSIONS: LEARNING TONAL GRAMMAR IN MUSIC

6.1. Separable Similarity and Emissions

In the HDP-HMM-LT model, we have a defined set of states with locations ℓ_j , $j = 1, \dots, J$. In Chapters 4 and 5, the similarities depended on the same parameters that determined the emission distributions, which we denoted by θ_j . In Chapter 4, each θ_j was a binary state vector, and the similarity $\phi_{jj'}$ was a decreasing function of the distance between those state vectors, while the emission distribution was a Gaussian centered at a linear function of θ_j . In Chapter 5, we generalized from binary state vectors to categorical state vectors, but similarity was still based on the distance between state vectors, and emission distributions were still Gaussians centered at a linear function of the state vector.

In this chapter, we suppose instead that ℓ_j consists of two separate and a priori independent parts: $\ell_j = (\theta_j, \eta_j)$, where the θ_j govern the emission distributions, and the similarities, $\phi_{jj'}$ depend on the η_j

Define

$$\phi_{jj'}(\eta_j, \eta_{j'}) = \exp\left(-\frac{\lambda}{2}\Delta_{jj'}^2\right)$$

where $\Delta_{jj'}$ is the Euclidean distance between η_j and $\eta_{j'}$; that is,

$$\Delta_{jj'}^2 = \sum_d (\eta_{jd} - \eta_{j'd})^2$$

6.2. A Hamiltionian Monte Carlo step to sample η

Since the η_j are real-valued vectors, we can sample them jointly using Hamiltonian Monte Carlo (HMC, also known as “Hybrid Monte Carlo”; Duane et al. (1987), see

also Neal et al. (2011)). HMC is a variation on the Metropolis-Hastings MCMC algorithm which is designed to more efficiently explore a high-dimensional continuous distribution by adopting a proposal distribution which is based on the evolution of Hamiltonian dynamics in a physical system. The position of the particle in the system represents the current state of the Markov chain, the potential energy of the particle is the negative log of the target distribution, and an auxiliary “momentum” variable is introduced, representing the kinetic energy of the system. The Markov chain evolves by computing a discrete approximation of an update to the position and momentum variables, and then computing the standard Metropolis-Hastings acceptance probability.

In order to carry out HMC in the context of the HaMMLLeT model with latent continuous state variables given by η_j , $j = 1, \dots, J$, we need the log likelihood and log prior for the η vector. Assume independent and isotropic Gaussian priors on each η_j , so we have

$$p(\eta_j) \propto \exp \left(-\frac{h_\eta}{2} \sum_d \eta_{jd}^2 \right),$$

where h_η is the prior precision which does not depend on d .

Then the log prior density, up to a constant, is

$$\log p(\eta_j) \propto -\frac{h_\eta}{2} \sum_d \eta_{jd}^2$$

The relevant log likelihood, as shown in Chapter 3 is the probability of the z and Q variables given the $\phi_{jj'}$. In particular, we have

$$L := p(\mathbf{z}, \mathbf{Q} | \phi) = \prod_j \prod_{j'} \phi_{jj'}^{n_{jj'}} (1 - \phi_{jj'})^{q_{jj'}}$$

and

$$\log L = \sum_j \sum_{j'} (n_{jj'} \log(\phi_{jj'}) + q_{jj'} \log(1 - \phi_{jj'}))$$

To do Hamiltonian Monte Carlo to sample from the conditional posterior of η

given \mathbf{z} and \mathbf{Q} , we need to compute the gradient of the log posterior, which is just the sum of the gradient of the log prior and the gradient of the log likelihood.

The j, d coordinate of the gradient of the log prior is simply

$$-2h_\eta\eta_{jd}$$

To get the j, d coordinate of the gradient of the log likelihood, we can apply the chain rule to terms as is convenient. In particular,

$$\frac{\partial L}{\partial \eta_{jd}} = \sum_j \sum_{j'} n_{jj'} \frac{\partial \log(\phi_{jj'})}{\partial \Delta_{jj'}^2} \frac{\partial \Delta_{jj'}^2}{\partial \eta_{jd}} + \sum_j \sum_{j'} q_{jj'} \frac{\partial \log(1 - \phi_{jj'})}{\partial (1 - \phi_{jj'})} \frac{\partial (1 - \phi_{jj'})}{\partial \Delta_{jj'}^2} \frac{\partial \Delta_{jj'}^2}{\partial \eta_{jd}}$$

We have the following components:

$$\begin{aligned} \frac{\partial \log(\phi_{jj'})}{\partial \Delta_{jj'}^2} &= -\frac{\lambda}{2} \\ \frac{\partial \Delta_{jj'}^2}{\partial \eta_{jd}} &= 2\Delta_{jj'd} I(j \neq j') \\ \frac{\partial \log(1 - \phi_{jj'})}{\partial (1 - \phi_{jj'})} &= \frac{1}{1 - \phi_{jj'}} \\ \frac{\partial (1 - \phi_{jj'})}{\partial \Delta_{jj'}^2} &= \frac{\lambda}{2} \phi_{jj'} \end{aligned}$$

which yields

$$\begin{aligned} \frac{\partial L}{\partial \eta_{jd}} &= -\lambda \sum_j \sum_{j'} n_{jj'} \Delta_{jj'd} \mathbb{I}(j \neq j') + \lambda \sum_j \sum_{j'} q_{jj'} \Delta_{jj'd} \frac{\phi_{jj'}}{1 - \phi_{jj'}} \mathbb{I}(j \neq j) \\ &= -\lambda \sum_{(j,j'): j \neq j'} \Delta_{jj'd} \left(n_{jj'} - q_{jj'} \frac{\phi_{jj'}}{1 - \phi_{jj'}} \right) \end{aligned}$$

$$p(\eta_j) \propto \exp \left(-\frac{h_\eta}{2} \sum_d \eta_{jd}^2 \right),$$

where h_η is the prior precision which does not depend on d .

6.3. Synthetic Data from an HMM with a Nearly Block Diagonal Transition Matrix

As a first check on this version of the model, we generated several datasets fixed state Hidden Markov models whose transition matrices were close to block-diagonal. Specifically, the state space consisted of 12 states in total which were grouped into 3 “superstates” of 4 states each. With probability 0.95, the chain transitioned to another state in the same superstate, and with probability 0.05 it transitioned to a different superstate. The distribution across the 4 same-superstate entries was drawn for each row from a symmetric Dirichlet, where the concentration parameter varied across datasets. The distribution for each row across the 8 other-superstate entries was drawn from a symmetric Dirichlet as well, where again, the concentration parameter was varied.

Each of the twelve states was associated with a mean in a two-dimensional observation space. However, unlike in the experiments discussed in Chapters 4 and 5, there was no relationship between proximity in transition space and proximity in emission space: the means for the twelve states were drawn independently from a $\mathcal{N}(0, \sigma^2 I)$ distribution, where σ^2 was varied across datasets.

At each time step, an observation was drawn from a bivariate $\mathcal{N}(\mu_j, I)$ distribution.

Since the HaMMLeT model infers a latent location for each state, and since the latent state locations impact the transition matrix by promoting transitions between nearby states and suppressing transitions between far away states, it is predicted that the HaMMLeT model will locate states in the same superstate near each other, and so we should expect to see a “clustering” of states into three distinct groups based on their η_j vectors. This clustering ability should allow HaMMLeT to more efficiently learn the correct transition matrix as compared to the HDP-HMM with no preference for local transitions.

To evaluate the performance of HaMMLeT versus the ordinary HDP-HMM, we

computed each models' calculated marginal log likelihood based on both the training data set used during inference, and on a held-out test set. At each Gibbs iteration, the marginal log likelihood was computed by fixing the transition matrix and emission parameters, and integrating out the state sequence using the forward message passing algorithm.

6.4. Discovering Chord Equivalence Classes in Tonal Music

A real-world test of the separable-similarity form of HaMMLeT comes from the music domain. We employed a dataset in which each observation consisted of four musical tones played concurrently, constituting a chord, and where the sequences between chords followed conventions for Western tonal music.

7. CONCLUSIONS AND FUTURE WORK

REFERENCES

- Baum, L. E. and Petrie, T. (1966). Statistical inference for probabilistic functions of finite state markov chains. *The annals of mathematical statistics*, 37(6):1554–1563.
- Beal, M. J., Ghahramani, Z., and Rasmussen, C. E. (2001). The infinite hidden markov model. In *Advances in neural information processing systems*, pages 577–584.
- Blei, D. M. and Jordan, M. I. (2006). Variational inference for Dirichlet process mixtures. *Bayesian Analysis*, 1(1):121–143.
- Box, G. E., Draper, N. R., et al. (1987). *Empirical model-building and response surfaces*, volume 424. Wiley New York.
- Duane, S., Kennedy, A. D., Pendleton, B. J., and Roweth, D. (1987). Hybrid monte carlo. *Physics letters B*, 195(2):216–222.
- Epanechnikov, V. A. (1969). Non-parametric estimation of a multivariate probability density. *Theory of Probability & Its Applications*, 14(1):153–158.
- Ferguson, T. S. (1973). A bayesian analysis of some nonparametric problems. *The annals of statistics*, pages 209–230.
- Fox, C. W. and Roberts, S. J. (2012). A tutorial on variational Bayesian inference. *Artificial intelligence review*, 38(2):85–95.
- Fox, E. B., Sudderth, E. B., Jordan, M. I., and Willsky, A. S. (2008). An hdp-hmm for systems with state persistence. In *Proceedings of the 25th international conference on Machine learning*, pages 312–319. ACM.
- Geman, S. and Geman, D. (1984). Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images. *IEEE Transactions on pattern analysis and machine intelligence*, 6:721–741.
- Gilks, W. R. and Wild, P. (1992). Adaptive rejection sampling for gibbs sampling. *Applied Statistics*, pages 337–348.
- Ishwaran, H. and Zarepour, M. (2000). Markov chain Monte Carlo in approximate Dirichlet and beta two-parameter process hierarchical models. *Biometrika*, 87(2):371–390.
- Johnson, M. J. and Willsky, A. S. (2013). Bayesian nonparametric hidden semi-markov models. *The Journal of Machine Learning Research*, 14(1):673–701.

- Kolter, J. Z. and Johnson, M. J. (2011). REDD: A public data set for energy disaggregation research. In *Workshop on Data Mining Applications in Sustainability (SIGKDD)*, San Diego, CA, volume 25, pages 59–62. Citeseer.
- Kurihara, K., Welling, M., and Teh, Y. W. (2007). Collapsed variational Dirichlet process mixture models. In *IJCAI*, volume 7, pages 2796–2801.
- Lindsay, B. G. (1995). Mixture models: theory, geometry and applications. In *NSF-CBMS regional conference series in probability and statistics*, pages i–163. JSTOR.
- McLachlan, G. and Peel, D. (2004). *Finite mixture models*. John Wiley & Sons.
- Neal, R. M. et al. (2011). MCMC using Hamiltonian dynamics. *Handbook of Markov Chain Monte Carlo*, 2:113–162.
- Paisley, J., Wang, C., and Blei, D. (2011). The discrete infinite logistic normal distribution. *arXiv preprint arXiv:1103.4789*.
- Parzen, E. (1962). On estimation of a probability density function and mode. *The annals of mathematical statistics*, 33(3):1065–1076.
- Rabiner, L. and Juang, B. (1986). An introduction to hidden markov models. *ieee assp magazine*, 3(1):4–16.
- Rasmussen, C. E. (2000). The infinite gaussian mixture model. In Solla, S., Leen, T., and Müller, K.-R., editors, *Advances in Neural Information Processing Systems*, volume 12, pages 554–560. MIT Press.
- Rosenblatt, M. et al. (1956). Remarks on some nonparametric estimates of a density function. *The Annals of Mathematical Statistics*, 27(3):832–837.
- Sethuraman, J. (1994). A constructive definition of Dirichlet processes. *Statistica Sinica*, 4:639–650.
- Teh, Y. W. (2011). Dirichlet process. In *Encyclopedia of machine learning*, pages 280–287. Springer.
- Teh, Y. W., Jordan, M. I., Beal, M. J., and Blei, D. M. (2006). Hierarchical Dirichlet processes. *Journal of the American Statistical Association*, 101(476).
- Wasserman, L. (2007). *All of nonparametric statistics*. Springer.